


CS3213: Foundations of Software Engineering

Course Summary

Bioblitz

- 751 observations
- 231 species
- 6 observers
- 142 identifiers



Sunda King Cobra



Wagler's Pit Viper



Oriental Whipsnake



Bioblitz

Orange-and-black Pear
Spider (354 observations)



By j_os_h_in

Saltwater Crocodile



Straw-headed Bulbul



Biodiversity Record: Long-tailed macaques mounting and grooming sambar deer

Manuel Rigger^{1*} & Jin Ting

¹ School of Computing, National University of Singapore, Singapore 117417
Email: rigger@nus.edu.sg (*corresponding author)

Recommended citation. Rigger M & Jin T (2026) Biodiversity Record: Long-tailed macaques mounting and grooming sambar deer. Nature in Singapore, 19: e2026024. DOI: 10.26107/NIS-2026-0024

Subjects: Long-tailed macaque, *Macaca fascicularis* (Mammalia: Primates: Cercopithecidae);
Sambar deer, *Rusa unicolor* (Mammalia: Artiodactyla: Cervidae).

Subjects identified by: Manuel Rigger and Jin Ting.

Location, date and time: Singapore Island, Mandai Track 15 (1°23'44.2"N 103°46'28.0"E); 25 January 2026; 1643–1648 hrs.

Habitat: Edge of secondary forest.

Observers: Manuel Rigger and Jin Ting.

Observation: A group of long-tailed macaques were observed interacting with a sambar. Up to four macaques at a time were sitting on the deer's back and grooming its dorsal area and ears (Fig. 1). The monkeys mounted the sambar over its hindquarters. The deer seemed calm and did not exhibit signs of distress toward the macaques. However, when the deer noticed the observers, it became alert, repeatedly tilting its head, causing the macaques around it to increase their distance. The macaques on the sambar's back dismounted when it started to move away. Two other sambar sighted in the vicinity were not interacting with monkeys.



Fig. 1 Four long-tailed macaques on a sambar, with at least two individuals grooming the dorsal region and ears of the deer (Screen grab from video recorded by: Manuel Rigger).

18th CITIZEN SCIENCE INFORMATION SESSION: RAFFLES' BANDED LANGUR SURVEYS IN SINGAPORE



Photo by Law Jia Bao

18 April 2026 Saturday 2-3pm

[A Zoom link will be sent to registered participants]

RAFFLES' BANDED LANGURS IN SINGAPORE

The Raffles' Banded Langur (*Presbytis femoralis*) is critically endangered in Singapore with a small population size and a restricted distribution. In order for conservation actions to be implemented, population estimate and the distribution of the Raffles' Banded Langurs need to be updated through field surveys. If you are interested to find out more on how to participate as a volunteer in the surveys for the Raffles' Banded Langurs in Singapore, please **register using the QR code below** to attend the online session.

ANDIE ANG

Andie leads the Species Grants and Impact at Mandai Nature and is the President of Jane Goodall Institute (Singapore). She chairs the Raffles' Banded Langur Working Group (RBL WG; <https://www.facebook.com/RBLWG>).

Organised by:



Supported by:



Register here





Agenda

- Course feedback
- Next week's presentations
- Exam mode
- Software engineering and AI
- Course repetition

Course Feedback (Blue)



Manuel Rigger AUTHOR | TEACHER

Created 27 Mar 11:17 | Posted 27 Mar 11:17

Submit course feedback (please!)

Hi everyone,

The whole teaching team would really appreciate it if you could take a few minutes to [submit your course feedback](#). As mentioned, we have put significant time and effort into trying to improve the course, and we will need your feedback to assess what worked well and what we still need to improve.

The key changes in the semester were the following:

- We overhauled the project component so that you can work on a hopefully exciting, meaningful, real-world project. We also changed the assignment format from a more open-ended agile one to a more concrete, but lightweight, plan-driven model.
- We converted the course to a blended-learning format to provide you with polished recordings and materials that you can study at home (based on the observation that most students did not come to class when we provided recordings of the lecture), while using the in-class time for polls, exercises, and demonstrations.
- We also re-created the tutorial contents, aiming for them to provide you with value by making them relevant for the exam or project.
- For the exam(s), we have tried to avoid questions that only test your memory, are very specific, or use the multiple-choice format.

We would also appreciate it if you could share your feedback on NUSMods after the end of the semester. This would allow next year's students to make a better-informed decision on whether to enroll in the course, assuming that some of the older feedback might no longer be applicable.

Your feedback really matters!

Reply



Student Feedback 2520

CS3213-FOUNDATIONS OF SOFTWARE ENGINEERING (LECTURE)

Live

43 Invited
0 Started
0 Responded
0 Opted Out

Evaluation ends on:
2026-04-24

0%
Response
Rate



Changes allowed until 2026-04-24





Follow-up Survey (Qualtrics)

Q2

You are invited to participate in an anonymous post-course survey about the course project. The goal of this survey is to evaluate the course design and to study how students' perspectives may have changed over the semester. This survey will take approximately 5 to 8 minutes to complete.

Your participation is voluntary, and choosing not to participate will not affect your course grade in any way. The survey does not collect identifying information. All responses will be analyzed and reported only in aggregated form. Aggregated results may be shared with the class and may be used in academic publications or institutional communications.

By continuing with the survey, you indicate your informed consent to participate.

Course Feedback (NUSMods)

Today

Timetable

Courses

Venues

Planner

Settings

Contribute

Whispers

NUSMODS

This new version of the prerequisite tree is being tested and may not be accurate. Viewers are encouraged to double check details with the prerequisite text above. To report bugs with the new tree, please post a bug report on GitHub (preferred) at [our repository](#) or send an email to bugs@nusmods.com.

Details

Prerequisites

Timetable

Reviews 6

Timetable

Semester 2

	1000	1100	1200	1300	1400	1500	1600	1700
MON					LEC [1] UT18			
TUE								
WED								
THU								
FRI	LAB [1] COM3-81-10 Weeks 2-13	LAB [2] COM3-81-10 Weeks 2-13	LAB [3] COM3-81-10 Weeks 2-13	LAB [4] COM3-81-10 Weeks 2-13	LAB [5] COM3-81-10 Weeks 2-13	LAB [6] COM3-81-10 Weeks 2-13		

Review and Discussion

6 Comments

Login

Join the discussion...

LOG IN WITH

OR SIGN UP WITH DISQUS

Name

Hi There!

We would like to encourage everyone who enjoyed using NUSMods to contribute back to the community by writing reviews for courses that you have taken before. Your efforts will go a long way in building up a vibrant and rich NUS community.

Please note:

1. Because the experience of each



Course Changes

- We **overhauled the project component** so that you can work on a hopefully exciting, meaningful, real-world project. We also changed the assignment format from a more open-ended agile one **to a more concrete, but lightweight, plan-driven model**.
- We converted the course to a **blended-learning format** to provide you with polished recordings and materials that you can study at home (based on the observation that most students did not come to class when we provided recordings of the lecture), while using the in-class time for **polls, exercises, and demonstrations**.
 - Including scavenger hunt and Bioblitz
- We also **re-created the tutorial contents**, aiming for them to provide you with value by making them relevant for the exam or project.
- For the **exam(s)**, we have tried to avoid questions that only test your memory, are very specific, or use the multiple-choice format.

GovTech (Chang Sau Sheong)



Deputy Chief Executive and Chief Technology Officer

GovTech Singapore · Full-time

Feb 2023 - Present · 3 yrs 2 mos

Singapore

GovTech is the government agency driving Singapore's Smart Nation initiative and public sector digital transformation. I lead GovTech teams to look at future technologies, set the technology standards and develop and deliver digital products for the Singapore government.

TEMASEK

Operating Partner, Digital

Temasek · Full-time

Jul 2022 - Feb 2023 · 8 mos

Singapore

I work with portfolio companies management teams to deliver value. During my time in Temasek, I was assigned to the AI Strategy and Solutions team and helped a few ventures, including Minden.AI's Yuu loyalty coalition, Alcadium, an AI technology company dedicated to creating and scaling AI solutions for enterprises, and others. I was also involved in AI technology funds, direct investments, and technical due diligence on potential investments.



Online Sharing Session

Would you attend an online sharing session by Chang Sau Sheong on April 15 (1:30pm to 3pm)?

1 Multiple choice *(Required)*



Would you attend a sharing session by Chang Sau Sheong on April 15 (1:30pm to 3pm)?

- ☐ Yes
- ☐ No
- ☐ Maybe

Submit

Final Submission

Assignment 6: Final Project Submission

CS3213 Foundations of Software Engineering (AY25/26 Sem 2)

Submission Deadline: **Fri, Apr 17, 2026, 11:59 pm**

-
- You must strictly comply with the noted deadline. No late submissions!
 - This is a **group** assignment. Acts of plagiarism are subjected to disciplinary action by the university. Please refer to <https://www.nus.edu.sg/celc/statements-and-e-resources-on-plagiarism/> for details on plagiarism and its associated penalties. *Note: it is sufficient if one member per group submits the assignment. One member may re-submit an updated version if needed.*
 - Please use appropriate tools to create your solutions (e.g., LibreOffice/Word/LaTeX for textual submissions, or **draw.io** for graphical solutions). Handwritten solutions are accepted only in exceptional cases and if they are very legible.
 - Please submit this PDF document via Canvas. In case of any discrepancies regarding the submission date, the date given in Canvas will count.
 - There are **20 points** to be scored for this assignment sheet. The worst score for any assignment sheet is 0 points.
-



Timeline

- By next Monday (May 11): prepare and give the presentation
- By next Friday (April 17): submit the report and artifact

This is the final deliverable for your project. Grading is based on the following weightage:

- **Live Demo** (8 points): Your ability to communicate what you built and demonstrate its key use cases and functionality purely through a live software demonstration. There will be no deliverable that you will need to submit on Canvas; rather, you will give a presentation in class.
- **Project report** (7 points): A concise report explaining your system architecture, design trade-offs, user value, and success metrics.
- **Project artifact** (5 points): The submitted software package (GitHub Release), including code quality and reproducibility.



Exam Mode

- 2 hours
- Closed book exam
- Double-sided A4 helpsheet is allowed



Exam Advice

- More is less!
- Think how to concisely answer the question before starting to write
- Stricter sentence limits than for the mid-term exam

6. Adhere strictly to the **sentence limits** provided. Points will be deducted for wordy answers or for exceeding the maximum sentence count.



Exam Advice

- The final exam will include questions about the first part of the semester that well
- Slightly more points for the contents of the second part of the semester



Exam Advice

- The second semester included more technical contents
- Prepare for more hands-on questions
- Examples (not exhaustive)
 - Specification-based testing (e.g., identifying partitions and boundary analysis)
 - Coverage computations (e.g., computing what coverage tests achieve, achieving 100% MC/DC coverage)
 - Applying mutation testing and computing the mutation score
 - Coming up with differential testing, metamorphic testing, and property-based testing schemes
 - Applying debugging strategies, delta-debugging, fault localization, and slicing



Exam Hints

- Similar format as mid-term exam
 - More open-text questions than the last year(s)
- No drawing of models
- No **in-depth questions** about the in-class research lecture
- Gave additional hints of possible questions during the semester

Previous Exams

- Previous exams are available only, but keep in mind that the contents changed over time as well as the different question format



Manuel Rigger AUTHOR | TEACHER

Created 30 Mar 11:05 | Posted 30 Mar 11:05



Final Exams from 2025 and 2024

For your reference and to help you prepare, I have attached the final exams from the previous two years and their solutions. Given the changes in lecture contents and improvement in exam format, this might be only partly helpful.

[CS3213 Exams 2024 and 2025.zip](#)

Reply





Software Engineering and AI



Image by ChatGPT

Course Exercise



Did you use LLMs for coding in your course project?

My Experience with Claude Code

- No errors. Here's a summary of all the changes made:

`pom.xml` - Upgraded `sqlite-jdbc` from 3.49.1.0 → 3.51.3.0

`SQLite3ExpressionGenerator.java` - Added to the `AnyFunction` enum:

- 28 math functions (3.35.0): `acos`, `acosh`, `asin`, `asinh`, `atan`, `atan2`, `atanh`, `ceil`, `ceiling`, `cos`, `cosh`, `degrees`, `exp`, `floor`, `ln`, `log`, `log2`, `log10`, `mod`, `pi`, `pow`, `power`, `radians`, `sin`, `sinh`, `sqrt`, `tan`, `tanh`, `trunc`
- `format()` alias for `printf` (3.38.0), `unixepoch()` (3.38.0)
- `unhex()` / `unhex(x, sep)` (3.41.0), `octet_length()` / `timediff()` (3.43.0)
- `concat()` / `concat_ws()` (3.44.0)
- Previously missing: `json_replace()`, `json_set()`
- 9 JSONB scalar functions (3.45.0): `jsonb`, `jsonb_array`, `jsonb_extract`, `jsonb_insert`, `jsonb_object`, `jsonb_patch`, `jsonb_remove`, `jsonb_replace`, `jsonb_set`
- `json_pretty()` (3.46.0), `unistr()` / `unistr_quote()` (3.50.0)
- Fixed `json_insert`, `json_object`, and all new odd/even-arg JSON functions to always generate the correct argument parity via `generateArguments` overrides

`SQLite3Expression.java` - Added:



- Binary operators `->` and `->>` (JSON extraction, 3.38.0)
- `IS DISTINCT FROM` and `IS NOT DISTINCT FROM` comparison operators (3.39.0)

`SQLite3Aggregate.java` - Added `STRING_AGG` aggregate function (3.44.0), with the generator updated to always pass 2 args for it

`SQLite3Errors.java` - Removed `json_insert()` needs an odd number of arguments and `json_object()` requires an even number of arguments since the generators now guarantee correct parity

* Brewed for 19m 12s

Course Exercise



How do you use LLMs in your coding workflow?

<https://pollev.com/manuelrigger>

Andrej Karpathy on Change



Andrej Karpathy ✓

@karpathy




It is hard to communicate how much programming has changed due to AI in the last 2 months: not gradually and over time in the "progress as usual" way, but specifically this last December. There are a number of asterisks but imo coding agents basically didn't work before December and basically work since - the models have significantly higher quality, long-term coherence and tenacity and they can power through large and long tasks, well past enough that it is extremely disruptive to the default programming workflow.

<https://x.com/karpathy/status/2026731645169185220>

Cursor: Agents Become More Common

Agents become mainstream and then the default

1.5x Agent requests per Tab accept



Michael Truell
@mntruell

The third era of AI software development

273

1.6K

9.8K



7.1M



When we started building Cursor a few years ago, most code was written one keystroke at a time. Tab autocomplete changed that and opened the first era of AI-assisted coding.

Then agents arrived, and developers shifted to directing agents through synchronous prompt-and-response loops. That was the second era. Now a third era is arriving. It is defined by agents that can tackle larger tasks independently, over longer timescales, with less human direction.

Course Exercise




What are the risks and challenges?

AI Tools: Risks (Meta)



Nothing humbles you like telling your OpenClaw "confirm before acting" and watching it speedrun deleting your inbox. I couldn't stop it from my phone. I had to RUN to my Mac mini like I was defusing a bomb.



11:25 AM · Feb 23, 2026 · 10M Views





Relevant ▾

View quotes >



IDEs/agent swarms/fallability. Both the "no need for IDE anymore" hype and the "agent swarm" hype is imo too much for right now. The models definitely still make mistakes and if you have any code you actually care about I would watch them like a hawk, in a nice large IDE on the side. The mistakes have changed a lot - they are not simple syntax errors anymore, they are subtle conceptual errors that a slightly sloppy, hasty junior dev might do. The most common category is that the models make wrong assumptions on your behalf and just run along with them without checking. They also don't manage their confusion, they don't seek clarifications, they don't surface inconsistencies, they don't present tradeoffs, they don't push back when they should, and they are still a little too sycophantic. Things get better in plan mode, but there is some need for a lightweight inline plan mode. They also really like to overcomplicate code and APIs, they bloat abstractions, they don't clean up dead code after themselves, etc. They will implement an inefficient, bloated, brittle construction over 1000 lines of code and it's up to you to be like "umm couldn't you just do this instead?" and they will be like "of course!" and immediately cut it down to 100 lines. They still sometimes change/remove comments and code they don't like or don't sufficiently understand as side effects, even if it is orthogonal to the task at hand. All of this happens despite a few simple attempts to fix it via instructions in CLAUDE . md. Despite all these issues, it is still a net huge improvement and it's very difficult to imagine going back to manual coding. TLDR everyone has their developing flow, my current is a small few CC sessions on the left in ghostty windows/tabs and an IDE on the right for viewing the code + manual edits.

Course Exercise



For what other tasks in SE can we use LLMs?

SE Tasks (until January 2024)

Large Language Models for Software Engineering: A Systematic Literature Review

XINYI HOU and YANJIE ZHAO, Huazhong University of Science and Technology, Wuhan, China

YUE LIU, Monash University, Melbourne, Australia

ZHOU YANG, Singapore Management University, Singapore

KAILONG WANG, Huazhong University of Science and Technology, Wuhan, China

LI LI, Beihang University, Beijing, China

XIAPU LUO, The Hong Kong Polytechnic University, Hong Kong, China

DAVID LO, Singapore Management University, Singapore

JOHN GRUNDY, Monash University, Melbourne, Australia

HAOYU WANG, Huazhong University of Science and Technology, Wuhan, China

Large Language Models (LLMs) have significantly impacted numerous domains, including Software Engineering (SE). Many recent publications have explored LLMs applied to various SE tasks. Nevertheless, a comprehensive understanding of the application, effects, and possible limitations of LLMs on SE is still in its early stages. To bridge this gap, we conducted a Systematic Literature Review (SLR) on LLM4SE, with a particular focus on understanding how LLMs can be exploited to optimize processes and outcomes. We selected and analyzed 395 research articles from January 2017 to January 2024 to answer four key Research Questions (RQs). In RQ1, we categorize different LLMs that have been employed in SE tasks, characterizing their distinctive features and uses. In RQ2, we analyze the methods used in data collection, pre-processing,

<https://dl.acm.org/doi/epdf/10.1145/3695988>

SE Tasks (until January 2024)

Requirements engineering	Anaphoric ambiguity treatment (4) Requirements analysis and evaluation (2) Coreference detection (1) Specification formalization (1) Use cases generation (1)	Requirements classification (4) Specification generation (2) Requirements elicitation (1) Traceability automation (1)	17
Software design	GUI retrieval (1) Software specification synthesis (1)	Rapid prototyping (1) System design (1)	4

SE Tasks (until January 2024)

Software development	Code generation (118) Code summarization (21) Code translation (12) API inference (5) API recommendation (5) Code representation (3) Method name generation (2) Agile story point estimation (1) API documentation smells (1) Data analysis (1) Control flow graph generation (1) Instruction generation (1) Others (14)	Code completion (22) Code search (12) Code understanding (8) Program synthesis (6) Code editing (5) Code comment generation (2) Code recommendation (2) API documentation augment (1) API entity and relation extraction (1) Fuzz driver generation (1) Identifier normalization (1) Type inference (1)	247
----------------------	--	--	-----

SE Tasks (until January 2024)

Software quality assurance	Vulnerability detection (18) Bug localization (5) Testing automation (4) Defect detection (2) Static analysis (2) Compiler fuzzing (1) Invariant prediction (1) Mobile app crash detection (1) Test prediction (1)	Test generation (17) Verification (5) Fault localization (3) GUI testing (2) Binary taint analysis (1) Decompilation (1) Malicious code localization (1) Resource leak detection (1)	66
----------------------------	--	---	----


SE Tasks (until January 2024)

Software maintenance	Program repair (35)	Code clone detection (8)	99
	Code review (7)	Debugging (4)	
	Bug reproduction (3)	Review/commit/code classification (3)	
	Duplicate bug report detection (3)	Logging (3)	
	Log parsing (3)	Code revision (2)	
	Sentiment analysis (3)	Vulnerability repair (2)	
	API misuses repair (1)	Bug prediction (1)	
	Bug triage (1)	Code coverage prediction (1)	
	Code review explained (1)	Code-Review defects repair (1)	
	Crash bug repair (1)	Dockerfile Repair (1)	
	Incivility detection (1)	Patch correctness prediction (1)	
	Patch detection (1)	Program merge conflicts repair (1)	
	Rename Refactoring (1)	Tag recommendation (1)	
	Technical debt payback (1)	Traceability recovery (1)	
	Web test repair (1)	Type error repair (1)	
	Others (5)		
Software management	Effort estimation (2)	Software tool configuration (1)	3

What Makes a Great Software Engineer?


Number of tokens
burned?

Number of concurrent
Claude Code windows?




What Makes a Great Software Engineer?

If your \$500K engineer isn't burning \$250K in tokens, something is wrong. - Jensen Huang



Course Exercise: Great Engineer



What Makes a Great Software Engineer?



<https://pollev.com/manuelrigger>

What Makes a Great Software Engineer?

2015 IEEE/ACM 37th IEEE International Conference on Software Engineering

What Makes a Great Software Engineer?

Paul Luo Li^{*,} Amy J. Ko^{*}, Jiamin Zhu⁺

Microsoft⁺
Seattle, WA
{pal,jiaminz}@microsoft.com

The Information School^{*}
University of Washington
ajko@uw.edu

Abstract—Good software engineers are essential to the creation of good software. However, most of what we know about software-engineering expertise are vague stereotypes, such as ‘excellent communicators’ and ‘great teammates’. The lack of specificity in our understanding hinders researchers from reasoning about them, employers from identifying them, and young engineers from becoming them. Our understanding also lacks breadth: what are all the distinguishing attributes of great engineers (technical expertise and beyond)? We took a first step in addressing these gaps by interviewing 59 experienced engineers across 13 divisions at Microsoft, uncovering 53 attributes of great engineers. We explain the attributes and examine how the most salient of these impact projects and teams. We discuss implications of this knowledge on research and the hiring and training of engineers.

Index Terms—Software engineers, expertise, teamwork

I. INTRODUCTION

Software engineering research has considered a vast number of factors that affect project outcomes, from process and tools, to programming languages and requirement elicitation. We rarely give consideration, however, to one of the most fundamental components of software engineering: the engineers themselves. Specifically, what makes a software engineer great? This basic question is at the foundation of nearly every part of our world’s rapidly growing software ecosystem: employers

In this study, we sought to remedy the lack of specificity, breadth, and rigor in prior work by investigating the following about *software* engineers:

- What do expert software engineers think are attributes of great software engineers?
- Why are these attributes important for the engineering of software?
- How do these attributes relate to each other?

To answer these questions, we performed 59 semi-structured interviews, spanning 13 Microsoft divisions, including several interviews with architect-level engineers with over 25 years of experience. The contribution of this effort is a thorough, specific, and contextual understanding of software engineering expertise, as viewed by expert software engineers.


In the rest of this paper, we detail our current understanding of software engineering expertise. We then discuss our interview and analysis methodology, the attributes we discovered, and the implications of this knowledge for software engineering research, practice, and training.

II. RELATED WORK


Much of our knowledge of software engineering expertise come from studying *new* engineers rather than experienced ones.

Face-to-face interview study at Microsoft to capture the opinion of experienced software engineers (in 2015)

What Makes a Great Software Engineer?



What Makes a Great Software Engineer?





Personal Characteristics: Improving

- **Improving:** not satisfied with the status quo and constantly looking to improve themselves, their product, and/or their surroundings
- Software engineering is a rapidly changing field

Personal Characteristics: Improving

*“Computer technology, compared to other sciences or technology, it's pretty young. Every year there's some new technology, new ideas. **If you are only satisfied with things you already learned, then you probably find out in a few years, you're out of date...** good software engineer [sic], he keep investigate, investment. [sic]”*
-SDE2, Corp Dev.

“... Always looking to do something better, always looking for the next thing, studying about the newer thing... to do things better.”
-Senior Dev Lead, Xbox

Personal Characteristics: Passionate


- **Passionate:** intrinsically interested in the area they are working in, and not just for extrinsic rewards such as money

*“You can't be a great engineer and not enjoy what you're doing... **not just to get a paycheck.**”*

-Principal SDE, Xbox

*“I think that there are people who are great software engineers who are in the wrong place and **aren't motivated and they end up not performing well.**” - Principal Dev Lead, Dynamics*

What Makes a Great Software Engineer?





Decision Making: People Knowledge

- **Knowledgeable about people and the organization:**
informed about the people around them: their responsibilities (i.e. organizational structure), their knowledge, and their tendencies

Decision Making: People Knowledge

*“Make sure that you **are aware of that big picture**, you know where you fit in and **how you interact with everyone else** to optimize what you are doing.”*

- Principal Dev Lead, Ad Platform

*“You have to understand people so **that you can influence or impact them...** You have to do that both down and up and out.”*

-Principal Dev Lead, Phone


Decision Making: Sees the Forest and the Trees

- **Sees the Forest and the Trees:** considering situations at multiple levels of abstraction, including technical details, industry trends, company vision, and customer/business needs

*“...Both a very, **very narrow extremely technical** prospective on his code, but also know where **it fits in with the bigger picture**, and to be aware of how it affects even our **major external customers, and the company vision.**”*

-Principal SDE, Windows

What Makes a Great Software Engineer?





Team Mates: Creates Shared Context

- **Creates Shared Context:** molding another person's understanding of a situation, was the most important aspect of “communicating effectively”
- Engineering involves many people, getting everyone to have a shared understanding was seen as essential to success



Team Mates: Creates Shared Context

*“One person can only accomplish so much so you've always got to be **working as part of the bigger group.** People **who can't communicate are only going to be sort of so-so effective...**”*
-Principal Dev Lead, Corp Dev

*“Most compellingly relate the value of that abstraction as it goes to non-abstract to very abstract to each person... **empathize with your audience...** get them to get it.”*
-SDE2, Windows


Team Mates: Creates a Safe Haven

- **Creates a Safe Haven:** creating a safe setting where engineers can learn and improve from mistakes and situations without negative consequences

“If you learn something from a failure, that’s a wonderful sort of thing... [but not] If you’re afraid of getting smacked upside the head... encourage the people to experiment, possibly succeed, possibly fail”.

-Senior SDE, Office

What Makes a Great Software Engineer?





Software Product: Elegant

- **Elegant:** simple and intuitive (i.e. not complex) software/designs that others can understand
- Among all software-related attributes, elegance was the most revered
- Complex solutions increase the **likelihood of bugs** and **increased maintenance costs** (if problems were fixable at all)
- **Evolving** the software is more costly when the design is brittle to change

Software Product: Elegant

*Never complicate any things... **when you simplify things it becomes easier for you to maintain**, going forward for customers... You get **lesser number of issues** reported by a customer.”*
-Senior Dev Lead, Dynamics

*“They are simple... very easy to understand in a sense that it’s very simple. **Doing something well and in a very simple way is very very hard.**”*
-Principal Dev Lead, Bing

Software Product: Anticipates Needs

- **Anticipates Needs:** accommodating possible future requirements not known at the time when the software was initially produced

“Think where you're going to get into trouble potentially... **What are we ultimately trying to do, and what can I do today that will save me time over the lifetime of the system?**”



-Technical Fellow



Conclusion

- Even before LLMs, programming skills were not what distinguished software engineers
- Many of the attributes relate to interaction with other humans (e.g., team mates, customers, the wider company, market, ...)

Course Exercise



What are the skills we should focus on?

Andrej Karpathy on Skills

- **Problem decomposition** (telling AI what to do clearly)
- **System thinking** (how pieces fit together)
- **Taste / judgment** (knowing if AI output is good)
- **Debugging & verification** (AI makes mistakes!)

As a result, programming is becoming unrecognizable. You're not typing computer code into an editor like the way things were since computers were invented, that era is over. You're spinning up AI agents, giving them tasks *in English* and managing and reviewing their work in parallel. The biggest prize is in figuring out how you can keep ascending the layers of abstraction to set up long-running orchestrator Claws with all of the right tools, memory and instructions that productively manage multiple parallel Code instances for you. The leverage achievable via top tier "agentic engineering" feels very high right now.

It's not perfect, it needs high-level direction, judgement, taste, oversight, iteration and hints and ideas. It works a lot better in some scenarios than others (e.g. especially for tasks that are well-specified and where you can verify/test functionality). The key is to build intuition to decompose the task just right to hand off the parts that work and help out around the edges. But imo, this is nowhere near "business as usual" time in software.

(summarized by ChatGPT)

<https://x.com/karpathy/status/2026731645169185220>

Gergely Orosz of The Pragmatic Engineer



More valuable:

- Being curious and experimenting with new tools/approaches
- Being interested in the business as well (not just the technology)
- Strong understanding of the fundamentals



Confidential - © 2020, SonarSource Sàrl

SonarSummit

125



Software Engineering and AI Skills

- Coding agents amplify engineering skills




eric zakariasson
@ericzakariasson



turns out, senior engineers accept more agent output than juniors. this is because:

- they write higher-signal prompts with tighter spec and minimal ambiguity
- they decompose work into agent-compatible units
- they have stronger priors for correctness, making review faster and more accurate
- juniors generate plenty but lack the verification heuristics to confidently greenlight output

shows that coding agents amplify existing engineering skill, not replace it



Summary by ChatGPT (Across Voices in AI)

✦ Common themes (what *actually stays important*)

Across all these people, there's a strong consensus:

🔑 1. Problem framing > coding

- Writing syntax matters less
- Defining the problem clearly matters more

🔑 2. Debugging & verification become core

- AI generates code → you *validate correctness*
- Testing skills become *more important*, not less

🔑 3. System design still king

- AI can write functions
- Humans design *architecture & tradeoffs*


🔑 4. Fundamentals don't go away

- Data structures, algorithms, performance
- Needed to *judge AI outputs*

🔑 5. Communication (with humans *and* AI)



- Prompting = structured thinking
- Explaining intent = critical skill

AI and Software Engineering



(Human-centric) software development processes

Course Exercise



Do you remember the course contents?