

CS4231  
Parallel and Distributed Algorithms

Lecture 1  
Instructor: YU Haifeng

## A bit of self-introduction...

- <https://www.comp.nus.edu.sg/~yuhf>

## References

- Some of the materials in the slides of this course are adapted from
  - Various textbooks (such as “Distributed Algorithms: An Intuitive Approach” by Wan Fokkink, 2018, 2nd edition)
  - Various recent research papers to keep the materials **up-to-date**
  - Some materials and homework problems from “Concurrent and Distributed Computing in Java” by Vijay Garg

# Module Overview

- Module homepage ready on Canvas
  - Check module homepage often!
- Polls on <https://www.polleverywhere.com/home> and [PollEv.com/haifengyu29](https://PollEv.com/haifengyu29)
- Prerequisite:
  - CS3230 Design and Analysis of Algorithms or CS3210 Parallel Computing --- CS3230 is more important
  - Can NOT be waived – will not honor request for waiver
- What is the module about:
  - Designing parallel/distributed algorithms, and proving their correctness/properties
  - This is a theory module, involving mostly with proofs and theorems

## Is This Module Hard?

- This is an **elective module** catering to students with strong interests and background in the subject
- While having the same theoretical nature as the **compulsory module CS3230**, this module will be much harder than CS3230
  - Analogy:
    - English modules taken by English major students
    - vs.
    - English modules taken by non-English-major students
  - This module will be taught in very different ways from CS3230

## Is This Module Hard?

- This is a 4000-level module:
  - Not recommended for 1<sup>st</sup>, 2<sup>nd</sup>, or 3<sup>rd</sup> year students – but if you think you can handle this challenging module, you can still take it but you are encouraged to chat with me before enrolling
- Overall, this is designed to be a challenging module
  - If you did not enjoy CS3230, you will not enjoy this module
  - If you enjoyed CS3230, you may still not enjoy this module if your formal/mathematical skill and abstract thinking ability are not strong
  - In previous offerings, it was not unusual that some students who did well in CS3230 had difficulty handling this module

## Why This Module Is Important

- This module is one of the most theoretical modules in SoC
  - Theoretical = impractical = useless?
  - No. *Actually it directly relates to your career...*
- Reason #1: Deep understanding of distributed algorithms **distinguishes you from others**
  - (Unfortunately) Everyone's career has a lot to do with competing with other people
  - Good programming skills ensure that you don't lose to others; deep understanding helps you to win

## Why This Module Is Important

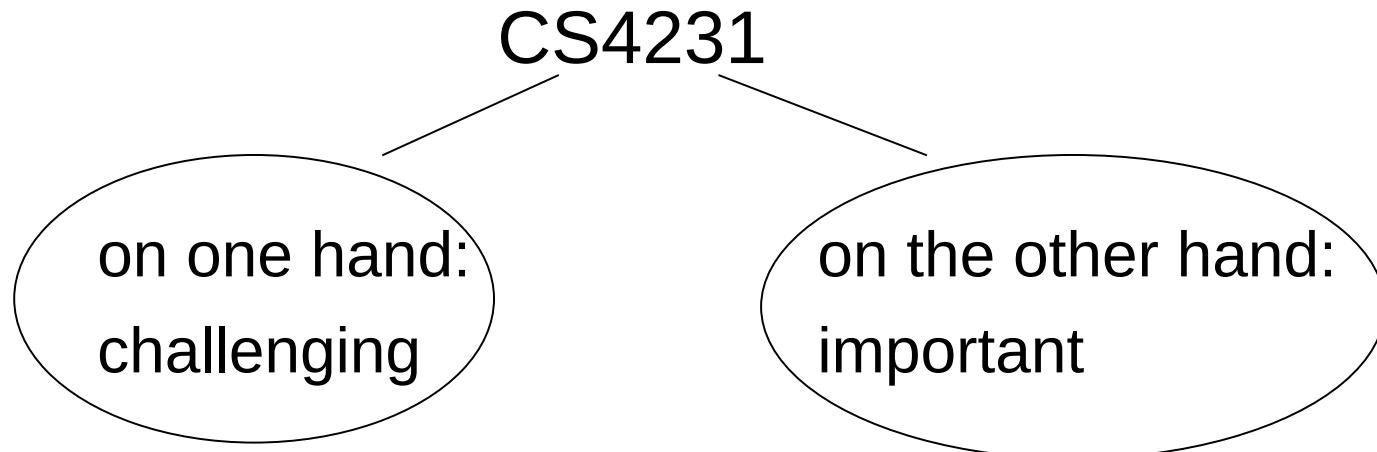
- Reason #2: Understanding foundational concepts **enables you to learn new material in the future**
  - Giving you some meat vs. giving you a hunting gun
  - Computer science is particularly fast-changing – don't expect what you learn today can be directly applicable 10 years later
- Reason #3: Foundational concepts are harder to grasp
  - This can be **your only opportunity** to learn these concepts in your whole career

## An Example

- Boss asks you to solve the following simple problem
  - Two nodes A and B, each has a starting value of 0 or 1
  - Each needs to output a single value 0 or 1
  - They can communicate but messages can be lost
  - Goal: A and B should output the same value. Specifically,
    - If A and B both start with 0, they should both output 0.
    - If A and B both start with 1 and if no messages are lost, they should both output 1.
    - If A and B both start with 1 and if some message is lost, they should **either** both output 1 **or** both output 0.
    - If A and B start with different values, they should **either** both output 1 **or** both output 0.

Quick Poll <https://pollev.com/haifengyu229>

To be, or not to be, that is the question



The decision is up to you...

## What students in previous years say

- “A very difficult module and I found it very hard to understand”
- “An interesting module that touch on parallel algorithms and distributed system, but concepts are a little too difficult sometimes”
- “The concepts are interesting, but really difficult to digest.”
- “Better to prove the algorithm in a more straight forward way. Too many lemma used, which make student hard to follow when using the lemma.”

## Relation with CS3211

- CS3211
  - PARALLEL AND CONCURRENT PROGRAMMING
- CS3211 is more programming oriented
- CS4231 is more about designing algorithms and proving their correctness/properties
- Overlap is small

## Teaching Format

- Wednesday 6:30pm-9:30pm (not 6:30pm-8:30pm) in every lecturing week: Lecturing + tutorial (break in the middle)
  - Face-to-face
- By NUS university guideline
  1. One-hour lesson: 45-minute teaching + concludes 15 minutes before the end of the hour
  2. Two-hour lesson: 90-minute teaching + 5-minute break in the middle + concludes 25 minutes before the end of the second hour
  3. Three-hour lesson: 135-minute teaching + 20-minute break in the middle (potentially split into two smaller breaks) + concludes 25 minutes before the end of the third hour
- You should read the materials to be covered **before each lecture** – Otherwise you won't follow

## Course Format

- Office hours:
  - Monday 4:00pm to 5:00pm every lecturing week in my office COM2-04-25
  - You are also welcome to approach me at other times or email me
- Do not email me via Canvas – please directly email [haifeng@comp.nus.edu  
u.sg](mailto:haifeng@comp.nus.edu.sg) to get faster response
- **Weekly homework – Use of AI not allowed**
- No systematic programming homework/exercise
  - To avoid overlapping with CS3211
  - To avoid excessive workload in this module
  - But you are still encouraged to implement algorithms learned

## Teaching Materials

- All materials taught in class and all materials in lecture notes are **compulsory** materials
  - These are kept up-to-date
- Why no textbooks?
  - Existing textbooks on the subject can be slightly old (as compared to the lecture notes)
  - Theoretical and foundational materials do not actually change so often, so people do not write a lot of new textbooks on such materials...
  - Some similar courses in other universities do not use textbooks either.

## Grading Policy

- 30% mid-term exam and 70% final exam
  - Both exams cover whatever have been taught by the time of the exam
  - Mid-term: Closed book, face-to-face
  - Final: Open book, face-to-face
- Homeworks do not directly contribute to final score
  - BUT some exam questions will be variants of homework questions
- Cheating ABSOLUTELY not tolerated and will be reported

## Mid-term Exam Date

- 19:00pm-20:00pm on Wed 04 Mar 2026 (students must arrive by 18:30pm)
  - Venue will be announce later
  - If venue unavailable, then it will be 19:00pm-20:00pm on Wed 11 Mar 2026
- Same policy as the final exam:
  - Not showing up for the mid-term exam = zero mark
  - Showing up later for the mid-term exam = less time to work on the exam (no extra time will be given)

## You MUST physically attend the (face-to-face) mid-term exam

- **Being able to physically attend the (face-to-face) mid-term exam is prerequisite for taking this module.**
- If you feel you have trouble showing up for the mid-term exam, let me know now...

## Class Participation and Penalty for Non-participation

- The School and also myself encourage class participation
  - But hard to incorporate into assessment (fairness issues)
- My approach to encourage participation:
  - If you miss any important discussions/announcements I made during lecture – you pay the price yourself
  - There will be important things that are only discussed during lecture, and not via email or Canvas or other venues
- Examples of questions that I will not answer:
  - I was at a party last night so I didn't attend the lecture, did you say anything important in the lecture?

## Rescheduling of Class

- ~~18 Feb 2026 is a public holiday, so we need to reschedule the class on that day...~~
  - ~~To 2:30pm-5:30pm on Saturday 24 Jan 2026~~

## Material Covered Today and Next Week

- Today: “Mutual Exclusion”
  - No tutorial today
- Next week:
  - “Synchronization Primitives”
  - Read slides before you come to class next week

# Break



# Mutual Exclusion Problem

- Context: Shared memory systems
  - Multi-processor computers
  - Multi-threaded programs
- Shared variable  $x$ 
  - Initial value 0
- Program  $x = x + 1$

<i>process 0</i>	<i>process 1</i>
read $x$ into a register (value read: 0)	
increment the register (1)	
write value in register back to $x$ (1)	
	read $x$ into a register (value read: 1)
	increment the register (2)
	write value in register back to $x$ (2)

# Mutual Exclusion Problem

- Context: Shared memory systems
  - Multi-processor computers
  - Multi-threaded programs
- Shared variable  $x$ 
  - Initial value 0
- Program  $x = x+1$
- Quick Poll <https://pollev.com/haifengyu229>

<i>process 0</i>	<i>process 1</i>
read $x$ into a register (value read: 0)	
increment the register (1)	
	read $x$ into a register (value read: 0)
	increment the register (1)
write value in register back to $x$ (1)	
	write value in register back to $x$ (1)

# Critical Section

Critical Section  
(also called  
Critical Region)

{ RequestCS(int processId)  
    Read  $x$  into a register  
    Increment the register  
    Write the register value back to  $x$   
    ReleaseCS(int processId)

# Roadmap

- Software solutions
  - Unsuccessful attempts
  - Peterson's algorithm
  - Bakery algorithm
- Hardware solutions
  - Disabling interrupts to prevent context switch
  - Special machine-level instructions

## Attempt 1

```
Shared boolean variable openDoor;  
    //whether door is open  
RequestCS(int processId) {  
    while (openDoor == false) {};  
    openDoor = false; // close door behind me  
}  
  
ReleaseCS(int processId) {  
    openDoor = true; // open door to let other people in  
}
```

Both process may  
see openDoor as  
true

Violate mutual exclusion: Two processes in critical region

## Attempt 2

Shared boolean variable wantCS[0], wantCS[1] initialized to false

wantCS[0] = wantCS[1] = true

<p>Process 0</p> <pre>RequestCS(0) {     wantCS[0] = true;     while (wantCS[1] == true) {}; }  ReleaseCS(0) {     wantCS[0] = false; }</pre>	<p>Process 1</p> <pre>RequestCS(1) {     wantCS[1] = true;     while (wantCS[0] == true) {}; }  ReleaseCS(1) {     wantCS[1] = false; }</pre>
---	---

No progress: No one can enter critical region

## Attempt 3

Shared int turn initialized to 0

*Process 0*

```
RequestCS(0) {  
    while (turn == 1) {};  
}
```

```
ReleaseCS(0) {  
    turn = 1;  
}
```

*Process 1*

```
RequestCS(1) {  
    while (turn == 0) {};  
}
```

```
ReleaseCS(1) {  
    turn = 0;  
}
```

Starvation: Process 0 may never enter critical region again  
(There are other kinds of starvation...)

## Properties Needed

1. **Mutual exclusion**: No more than one process in the critical section
2. **Progress**: If one or more process wants to enter and if no one is in the critical section, then one of them can eventually enter the critical section

Intuitively, we want to ensure that “resource” is fully utilized
3. **No starvation**: If a process wants to enter, it eventually can always enter

No starvation implies Progress

  - Always need to consider the worst-case schedule.
  - If a process is in critical section, we always assume that it will eventually exit the critical section.

# Peterson's Algorithm

Shared bool wantCS[0] = false, bool wantCS[1] = false, int turn = 0;

*Process 0*

```
RequestCS(0) {  
    wantCS[0] = true;  
    turn = 1;  
    while (wantCS[1] == true &&  
          turn == 1) {};  
}
```

```
ReleaseCS(0) {  
    wantCS[0] = false;  
}
```

*Process 1*

```
RequestCS(1) {  
    wantCS[1] = true;  
    turn = 0;  
    while (wantCS[0] == true &&  
          turn == 0) {};  
}
```

```
ReleaseCS(1) {  
    wantCS[1] = false;  
}
```

# Correctness Proof for Peterson's Alg.

*Process 0*

```
RequestCS(0) {  
    wantCS[0] = true;  
    turn = 1;  
    while (wantCS[1] == true &&  
          turn == 1) {};  
}
```

*Process 1*

```
RequestCS(1) {  
    wantCS[1] = true;  
    turn = 0;  
    while (wantCS[0] == true &&  
          turn == 0) {};
```

**Mutual exclusion:** Proof by contradiction.

Case 1:  $\text{turn} == 0$  when P0 and P1 are both in critical section.

Then P0 executed “ $\text{turn} = 1$ ” before P1 executed “ $\text{turn} = 0$ ”.

Hence  $\text{wantCS}[0] == \text{false}$  as seen by P1.

But  $\text{wantCS}[0]$  set to true by Process 0.

Case 2:  $\text{turn} == 1$ . Symmetric – complete yourself...

## Correctness Proof for Peterson's Alg.

*Process 0*

```
RequestCS(0) {  
    wantCS[0] = true;  
    turn = 1;  
    while (wantCS[1] == true &&  
          turn == 1) {};  
}
```

*Process 1*

```
RequestCS(1) {  
    wantCS[1] = true;  
    turn = 0;  
    while (wantCS[0] == true &&  
          turn == 0) {};
```

**Progress:** Proof by contradiction and consider the value of turn when both P0 and P1 are waiting.

Case 1:  $\text{turn} == 0$ . Then P0 can enter.

Case 2:  $\text{turn} == 1$ . Symmetric – complete yourself...

# Correctness Proof for Peterson's Alg.

*Process 0*

```
RequestCS(0) {  
    wantCS[0] = true;  
    turn = 1;  
    while (wantCS[1] == true &&  
          turn == 1) {};  
}
```

*Process 1*

```
RequestCS(1) {  
    wantCS[1] = true;  
    turn = 0;  
    while (wantCS[0] == true &&  
          turn == 0) {};  
}  
ReleaseCS(1) {  
    wantCS[1] = false;  
}
```

**No starvation:** Proof by contradiction.

Case 1: If P0 waiting, then  $\text{wantCS}[1] = \text{true}$  and  $\text{turn} = 1$ .

P1 in critical region -- will exit and set  $\text{wantCS}[1]$  to false.

What if P1 wants to enter again immediately?

Case 2: P1 is waiting. Symmetric – complete yourself...

## Lamport's Bakery Algorithm

- For  $n$  processes
  - Get a number first
  - Get served when all people with lower number have been served
- Two shared arrays of  $n$  elements
  - boolean choosing[ $i$ ] = false; // process  $i$  is trying to get a number
  - int number[ $i$ ] = 0; // the number got by process  $i$ ;  
// “0” means process  $i$  not interested in being served

```
ReleaseCS(int myid) {  
    number[myid] = 0;  
}  
  
// a utility function  
boolean Smaller(int number1, int id1, int number2, int id2) {  
    if (number1 < number 2) return true;  
    if (number1 == number2) {  
        if (id1 < id2) return true; else return false;  
    }  
    if (number 1 > number2) return false;  
}
```

```
RequestCS(int myid) {  
    choosing[myid] = true;  
    for (int j = 0; j < n; j++)  
        tmp = number[ j ];  
        if (tmp > number[myid]) number[myid] = tmp;  
    number[myid]++;  
    choosing[myid] = false;  
  
    for (int j = 0; j < n; j++) {  
        while (choosing[ j ] == true);  
        while (number[ j ] != 0 &&  
            Smaller(number[ j ], j, number[myid], myid));  
    }  
}
```

get a  
number

wait for  
people  
ahead  
of me

```

choosing[myid] = true;
for (int j = 0; j < n; j++)
    tmp = number[j];
    if (tmp > number[myid])
        number[myid] = tmp;
number[myid]++;
choosing[myid] = false;

for (int j = 0; j < n; j++) {
    while (choosing[j] == true);
    while (number[j] != 0 &&
           Smaller(number[j], j,
           number[myid], myid));
}

```

- **Progress:** Proof by contradiction. Consider any set of processes that wants to enter the CS but no one can make progress. Each process is guaranteed to get a queue #. Let process  $i$  be the one with the smallest queue number. Consider where process  $i$  can be blocked:
  - Case 1:  
Process  $j$  will eventually set  $\text{choosing}[j]$  to false  
Process  $j$  will then block  
(otherwise there is progress already!)
  - Case 2: Impossible since process  $i$  has the smallest queue number
- **No starvation:** Can be similarly shown...work it out yourself...

Mutual exclusion: Suppose  $i$  and  $k$  both in critical section.

At T1, process  $i$  is here

```
choosing[myid] = true;  
for (int j = 0; j < n; j++)  
    tmp = number[j];  
    if (tmp > number[myid])  
        number[myid] = tmp;  
number[myid]++;  
choosing[myid] = false;  
  
for (int j = 0; j < n; j++) {  
    while (choosing[j] == true);  
    while (number[j] != 0 &&  
        Smaller(number[j], j,  
        number[myid], myid));  
}
```

At T1, process  $k$  is here

- W.l.o.g, assume  $\text{Smaller}(\text{number}[ i ], i, \text{number}[ k ], k)$  after they are in the critical sec
- Process  $k$  must see  $\text{number}[i] == 0$  at that time T1:  
We want to know where process  $i$  is at time T1.
  - Case 1: Process  $i$  has not executed  
“ $\text{tmp} = \text{number}[ k ]$ ”. Then eventually  $\text{number}[ i ] > \text{number}[ k ]$ . Impossible.
  - Case 2: Has executed “ $\text{tmp} = \text{number}[ k ]$ ”
    - Subcase 2.1: Process  $i$  has executed  
“ $\text{number}[myid]++;$ ”  
-- impossible since  $\text{number}[ i ] == 0$
    - Subcase 2.2: Has not executed “ $\text{number}[myid]++;$ ”  
-- This is the only possible case.

At T1, process  $i$  is here

```

choosing[myid] = true;
for (int j = 0; j < n; j++)
    tmp = number[j];
    if (tmp > number[myid])
        number[myid] = tmp;
number[myid]++;
choosing[myid] = false;
```

At T2, process  $k$  is here

```

for (int j = 0; j < n; j++) {
    while (choosing[j] == true);
    while (number[j] != 0 &&
           Smaller(number[j], j,
                   number[myid], myid));
}
```

At T1, process  $k$  is here

- Now continue and consider the time T2 when process  $k$  invoked “while (choosing[  $i$  ] = true);” and passed that statement.
- We want to see where process  $i$  is at T2. Since choosing[  $i$  ] = false, process  $i$  must either have finished choosing its queue number or have not started choosing:
  - Case 1: process  $i$  has finished choosing and has executed choosing[  $i$  ] = false; Impossible since  $T2 < T1$ .
  - Case 2: process  $i$  has not started choosing and has not executed choosing[  $i$  ] = true. But then number[  $i$  ] will be larger than number[  $k$  ]. Contradiction.

# Roadmap

- Software solutions
  - Unsuccessful attempts
  - Peterson's algorithm
  - Bakery algorithm
- Hardware solutions
  - Disabling interrupts to prevent context switch
  - Special machine-level instructions

# Disable Interrupts

Do not allow context switch here

process 1	process 2
read $x$ into a register (value read: 0)	
increment the register (1)	
	read $x$ into a register (value read: 0)
	increment the register (1)
write value in register back to $x$ (1)	
	write value in register back to $x$ (1)

# Special Machine-level Instructions

```
boolean TestAndSet(Boolean openDoor, boolean newValue) {  
    boolean tmp = openDoor.getValue();  
    openDoor.setValue(newValue);  
    return tmp;  
}
```

Executed  
atomically

---

shared Boolean variable openDoor initialized to true;  
RequestCS(process\_id) {  
 while (TestAndSet(openDoor, false) == false) {};  
}  
ReleaseCS(process\_id) { openDoor.setValue(true); }

# Summary

- Module overview & policy
- Mutual exclusion problem in shared-memory systems
- Software solutions
  - Unsuccessful attempts
  - Peterson's algorithm
  - Bakery algorithm
- Hardware solutions
  - Disabling interrupts to prevent context switch
  - Special machine-level instructions

## Homework Assignment (on this and next few slides)

- Devise a mutual exclusion algorithm for  $n$  processes by using Peterson's 2-mutual-exclusion algorithm **as a black-box**
- Construct two interleaving executions to show that each of the following modifications to Peterson's algorithm will make it incorrect:
  - A process sets the **turn** variable to itself instead of to the other process.
  - A process sets the **turn** variable before setting the **wantCS** variable.
- Construct an interleaving executions to show that the bakery algorithm becomes incorrect, if we do not use the **choosing** variable

- Does the following algorithm for 2 processes guarantee: Mutual exclusion? Progress? No starvation? For each of these 3 properties, either give a proof or construct a problematic scenario as counter-example.

Shared boolean wantCS[0] = false, boolean wantCS[1] = false, int turn = 1;

*Process 0*

```
RequestCS(0) {
    wantCS[0] = true;
    while (wantCS[1]) {
        if (turn == 1) {
            wantCS[0] = false;
            while (turn == 1) { ;}
            wantCS[0] = true;
        }
    }
    ReleaseCS(0) { turn = 1; wantCS[0] = false; }
```

*Process 1*

```
RequestCS(1) {
    wantCS[1] = true;
    while (wantCS[0]) {
        if (turn == 0) {
            wantCS[1] = false;
            while (turn == 0) { ;}
            wantCS[1] = true;
        }
    }
    ReleaseCS(1) { turn = 0; wantCS[1] = false; }
```

- Bring you **completed** homework to class next week