

CS4231
Parallel and Distributed Algorithms

Lecture 3
Instructor: YU Haifeng

Review of Last Lecture

- Why do we need synchronization primitives
 - Busy waiting waste CPU
 - But synchronization primitives need OS support
- Semaphore:
 - Using semaphore to solve dining philosophers problem
 - Avoiding deadlocks
- Monitor:
 - Easier to use than semaphores / more popular
 - Two kinds of monitors
 - Using monitor to solve producer-consumer and reader-writer problem

Today's Roadmap

- “Consistency Conditions”
- What is consistency? Why do we care about it?
- Sequential consistency
- Linearizability
- Consistency models for registers

Abstract Data Type

- Abstract data type: A piece of data with allowed operations on the data
 - Integer X, read(), write()
 - Integer X, increment()
 - Queue, enqueue(), dequeue()
 - Stack, push(), pop()
- We consider **shared** abstract data type
 - Can be accessed by multiple processes
 - Shared object as a shorthand

What Is Consistency?

- Consistency – a term with thousand definitions
- Definition in this course:
 - Consistency specifies what behavior is allowed when a shared object is accessed by multiple processes
 - When we say something is “consistent”, we mean it satisfies the specification (according to some given spec)
- Specification:
 - No right or wrong – anything can be a specification
- But to be useful, must:
 - Be sufficiently strong – otherwise the shared object cannot be used in a program
 - Can be implemented (efficiently) – otherwise remains a theory
 - Often a trade-off

In Lecture #1: Mutual Exclusion Problem with $x = x+1$

- Data type 1:
 - Integer x
 - `read()`, `write()`
- Data type 2:
 - Integer x
 - `increment()`
- We were using type 1 to implement type 2
 - By doing $x = x+1$;

<i>process 0</i>	<i>process 1</i>
read x into a register (value read: 0)	
increment the register (1)	we implicitly assumed the value read must be the value written – one possible def of consistency
write value in register back to x (1)	
	read x into a register (value read: 1)
	increment the register (2)
	write value in register back to x (2)

In Lecture #1: Mutual Exclusion Problem with $x = x+1$

- Data type 1:
 - Integer x
 - `read()`, `write()`
- Data type 2:
 - Integer x
 - `increment()`
- By saying the execution is not “good”, we implicitly assumed some consistency definition for Data type 2

<i>process 0</i>	<i>process 1</i>
read x into a register (value read: 0)	
increment the register (1)	
	read x into a register (value read: 0)
	increment the register (1)
write value in register back to x (1)	
	write value in register back to x (1)

Mutual Exclusion and Consistency

- Mutual exclusion is actually a way to ensure consistency (based on some consistency definition)
 - Thus we actually already dealt with consistency...

Sequential Consistency

- Sequential consistency arguably the most widely used consistency definition
 - Almost all **commercial databases** such as Oracle databases
 - Typical **multi-processors** and **multi-core CPUs**
- First defined by Lamport: “...the results ... is the same as if the operations of all the processors were executed in some sequential order, and the operations of each individual processor appear in this sequence in the order specified by the program”
 - Use sequential order as a comparison point
 - Focus on results only – that is what users care about (what if only allow sequential histories?)
 - Require that program order is preserved
- We need a lot of formalism to formalize the above...

Formalizing a Parallel Execution

- *Operation:* A single invocation/response pair of a single method of a single shared object by a process
 - e being an operation
 - proc(e): The invoking process
 - obj(e): The object
 - inv(e): Invocation **event** (*start time*)
 - resp(e): Reply **event** (*finish time*)
- Two invocation events are the same if invoker, invokee, parameters are the same –
 $\text{inv}(p, \text{read}, X)$
- Two response events are the same if invoker, invokee, response are the same –
 $\text{resp}(p, \text{read}, X, 1)$



wall clock time /
physical time /
real time

(Execution) History

- A *history H* is a sequence of invocations and responses ordered by wall clock time
 - For any invocation in H, the corresponding response is required to be in H
 - Each execution of a parallel system corresponds to a history, and vice versa.


inv(p, read, X) inv(q, write, X, 1) resp(p, read, X, 0) resp(q, write, X, OK)

Sequential History and Concurrent History

- A history H is *sequential* if
 - Any invocation is always *immediately* followed by its response
 - I.e. No interleaving
 - Otherwise called *concurrent*

Sequential:


$\text{inv}(p, \text{read}, X) \text{ resp}(p, \text{read}, X, 0) \text{ inv}(q, \text{write}, X, 1) \text{ resp}(q, \text{write}, X, \text{OK})$



The diagram shows two pairs of messages: (inv(p, read, X), resp(p, read, X, 0)) and (inv(q, write, X, 1), resp(q, write, X, OK)). Each pair is grouped together by a curly brace, indicating they happen sequentially.

concurrent:

$\text{inv}(p, \text{read}, X) \text{ inv}(q, \text{write}, X, 1) \text{ resp}(p, \text{read}, X, 0) \text{ resp}(q, \text{write}, X, \text{OK})$



The diagram shows four messages in an interleaved sequence: inv(p, read, X), inv(q, write, X, 1), resp(p, read, X, 0), and resp(q, write, X, OK). There are three curly braces: one spanning the first two messages, one spanning the last two messages, and one large brace at the bottom spanning all four messages, indicating they can happen in any order.

Sequential Legal History and Subhistory

- A sequential history H is *legal* if
 - All responses satisfies the *sequential semantics* of the data type
 - *Sequential semantics*: The semantics you would get if there is only one process accessing that date type.
 - It is possible for a sequential history not to be legal
- Process p 's *process subhistory* of H ($H \mid p$):
 - The subsequence of all events of p
 - Thus a process subhistory is always sequential
- Object o 's *object subhistory* of H ($H \mid o$):
 - The subsequence of all events of o


Equivalency and Process Order

- Two histories are *equivalent* if they have the exactly same set of events
 - Same events imply **all responses are the same**
 - Ordering of the events may be different
 - **User only cares about responses**
- *Process order* is a partial order among all events
 - For any two events of the same process, process order is the same as execution order
 - No other additional orderings



Sequential Consistency

- A history H is sequentially consistent if it is equivalent to some legal sequential history S that preserves process order
 - Use sequential history as a comparison point
 - Focus on results only – that is what users care about
 - Require that program order be preserved


Examples (initial value of x is 0)





$\sim\equiv$



$\sim\equiv$




Examples (initial value of x is 0)



Quick Poll <https://pollev.com/haifengyu229>


Motivation for Linearizability

- Sequential consistency sometimes not strong enough




More Formalism

- A history H uniquely induces the following “ $<$ ” partial order among operations
 - $o_1 < o_2$ iff response of o_1 appears in H before invocation of o_2
 - We call this *external order*
- What is the external order induced by a sequential history?



Linearizability

- Definition #1: The execution is equivalent to some execution such that each operation happens instantaneously at some point (called **linearization point**) between the invocation and response
 - Can be made rigorous – will be your homework




The above history is sequentially consistent but not linearizable

Linearizability

- Definition #1 (repeated): The execution is equivalent to some execution such that each operation happens instantaneously at some point between the invocation and response
- Definition #2: History H is linearizable if
 - It is equivalent to some legal sequential history S, and
 - S preserves the external order in H – namely, the partial order induced by H is a subset of the partial order induced by S
- Are the two definitions the same? This question will be your homework today...

Linearizability

- A linearizable history must be sequentially consistent
 - Linearizability is arguably the strongest form of consistency people have ever defined




Another interesting example

Local Property


- Linearizability is a **local property** in the sense that H is linearizable if and only if for any object x , $H \upharpoonright x$ is linearizable
- Sequential consistency is not a local property

Sequential Consistency is Not Local Property

x, y are initially empty queues




$H \mid x$ is sequentially consistent:



Similarly, $H \mid y$ is sequentially consistent as well

Sequential Consistency is Not Local Property

x, y are initially empty queues



But H is not sequentially consistent


Proving That Linearizability is a Local Property

- Definition #1 (repeated): The execution is equivalent to some execution such that each operation happens instantaneously at some point between the invocation and response
- The proof is trivial if we use the above definition (after formalizing it)
- But we'll prove the statement while using Definition #2

Proving That Linearizability is a Local Property

- Definition #2 (repeated): A history H is linearizable if
 1. It is equivalent to some legal sequential history S , and
 2. S preserves the external order in H
- Want to show H is linearizable if for any object x , $H \mid x$ is linearizable
- We construct a directed graph among all operations as following: A directed edge is created from o_1 to o_2 if
 - o_1 and o_2 is on the same obj x and o_1 is before o_2 when linearizing $H \mid x$ (we say $o_1 \rightarrow o_2$ due to obj), OR
 - $o_1 < o_2$ in external order (i.e., response of o_1 appears before invocation of o_2 in H) (we say $o_1 \rightarrow o_2$ due to H)

Proving That Linearizability is a Local Property



- Lemma: The resulting directed graph is acyclic (we prove this later)
- Any topological sorting of the above graph gives us a legal sequential history S, and the edges in the graph is a subset of the total order induced by S – Theorem proved

Linearizability definition:

1. H is equivalent to some **legal** sequential history S, and
2. S preserves the external order in H

Legal + Legal = Legal

- $H \mid x$ is linearizable, and let the linearization be O₁ O₂
- $H \mid y$ is linearizable, and let the linearization be O₃ O₄
- Suppose the topological sort be O₁ O₃ O₂ O₄
- O₁ O₂ is legal (for x), O₃ O₄ is legal (for y)
- Is S = O₁ O₃ O₂ O₄ legal (for both x and y)?
 - In particular, we are inserting O₃ between O₁ and O₂, will O₂'s behavior change due to some value read in O₃?
- All events in S are from H – **including all invocation events and their parameters**
 - So if in H, the program code updates y based on the value of x, the new value for y will already be captured in the parameters of the corresponding invocation event – we do not re-run the program code

Lemma: The resulting graph does not have a cycle


operations on obj1



operations on obj2




operations on obj3




- Prove by contradiction

Lemma: The resulting graph does not have a cycle




- Example: The red cycle is composed of
 - Edges due to obj3, followed by
 - Edges due to H, followed by
 - Edges due to obj2, followed by
 - Edges due to H, followed by

Lemma: The resulting graph does not have a cycle



- Impossible to have
 - Edges due to some object x, followed by
 - Edges due to some object y


Lemma: The resulting graph does not have a cycle



- To generalize, any cycle must be composed of:
 - Edges due to some object x, followed by
 - Edges due to H, followed by
 - Edges due to some object y, followed by
 - Edges due to H, followed by
 -


Lemma: The resulting graph does not have a cycle

edges due to x



$H \mid x$ is equivalent to a serial history S – S is a total order


edges due to H



the partial order induced by H is transitive


- To generalize, any cycle must be composed of:
 - Edges due to some object x → A single edge due to x
 - Edges due to H → A single edge due to H
 - Edges due to some object y → A single edge due to y
 - Edges due to H → A single edge due to H
 -

Lemma: The resulting graph does not have a cycle




- We must have a cycle in the form of
 - A single edge due to x , followed by
 - A single edge due to H , followed by
 - A single edge due to y , followed by
 - A single edge due to H , followed by
 -

Lemma: The resulting graph does not have a cycle



in H :



From the edge $D \rightarrow A$:

From the edge $A \rightarrow B$:

From the edge $B \rightarrow C$:


It is impossible to have $C \rightarrow D$ due to y

Consistency Definitions for Registers

- **Register** is a kind of abstract data type: A single value that can be read and written
- An (implementation of a) register is called **atomic** if the implementation always ensures **linearizability** of the history
- An (implementation of a) register is called **??** if the implementation always ensures **sequential consistency** of the history

Consistency Definitions for Registers

- An (implementation of a) register is called **regular** if
 - When a read does not **overlap with any write**, the read returns the value written by **one of the most recent writes**
 - When a read overlaps with one or more writes, the read returns the value written by **one of the most recent writes or the value written by one of the overlapping writes**
- Definition of most recent writes





An atomic register must
be regular

Regular implies Sequential Consistency ?

Sequential Consistency implies Regular ?

assuming the initial value of the register is 0



Consistency Definitions for Registers

- An (implementation of a) register is called **safe** if the implementation always ensures that
 - When a read does not overlap with any write, then it returns the value written by **one of the most recent writes**
 - When a read overlaps with one or more writes, it can return anything

A regular register hence must be safe

- Safe implies Sequential Consistency ? Yes/No
- Sequential Consistency implies Safe ? Yes/No



Quick Poll <https://pollev.com/haifengyu229>

Summary

- What is consistency? Why do we care about it?
- Sequential consistency
- Linearizability
 - Linearizability is a local property
- Consistency models for registers

Homework Assignment (on this and next few slides)

- Consider each of the following 2 histories, where S is a shared stack (initially empty):
 - If you believe the history is linearizable/Serializable, give the equivalent sequential history
 - If you believe it is not linearizable/Serializable, **prove** that no equivalent sequential history exists



- Prove that the two definitions of linearizability on Slide 21 are equivalent
- Prove that linearizability is a local property using the definition on slide 26 (you need to formalize that definition first)
- Bring your completed homework to class next week