

CS4231
Parallel and Distributed Algorithms

Lecture 4

Instructor: YU Haifeng

Review of Last Lecture

- What is consistency? Why do we care about it?
- Sequential consistency
- Linearizability
 - Linearizability is a local property
- Consistency models for registers

Today's Roadmap

- “Models and Clocks”
 - Goal: Define “time” in a distributed system
- Logical clock
- Vector clock
- Matrix clock

Assumptions

- Process can perform three kinds of atomic actions/events
 - Local computation
 - Send a single message to a single process
 - Receive a single message from a single process
 - No atomic broadcast
- Communication model
 - Point-to-point
 - Error-free, infinite buffer
 - Potentially out of order

Motivation

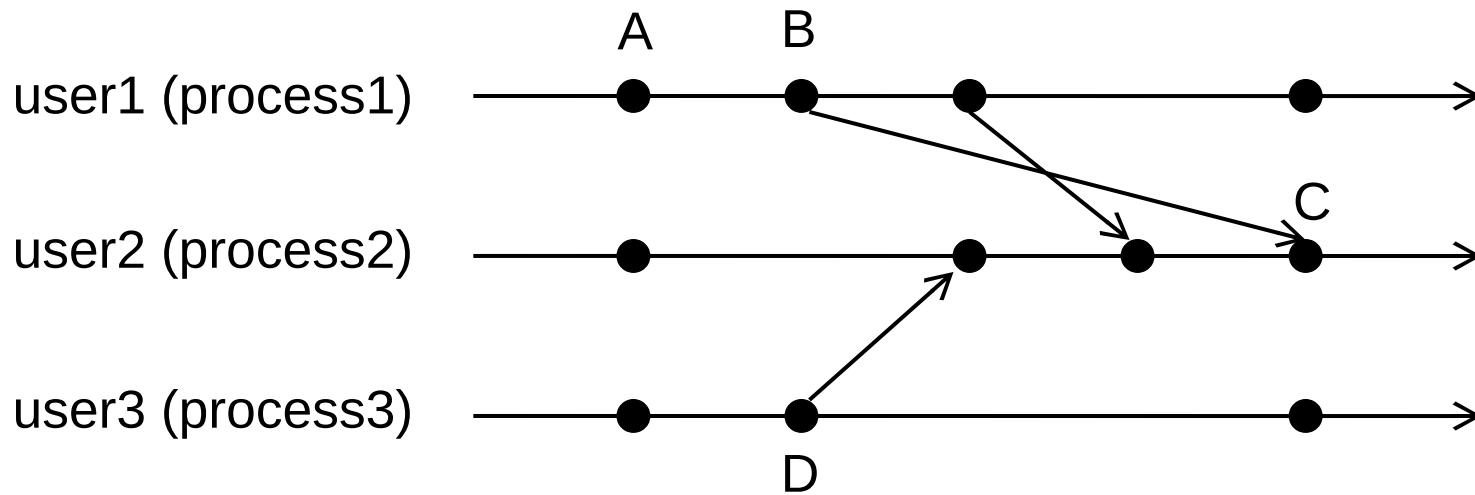
- Many protocols need to impose an ordering among events
 - Event A: You mom deposit some money into your bank account as your birthday gift
 - Event B: You use your ATM card to use the money to buy some stuff
 - Your bank needs to properly order the two events
- Physical clocks:
 - Seems to completely solve the problem
 - But what about theory of relativity?
 - Even without theory of relativity – efficiency problems
- How accurate is sufficient?
 - Clock error needs to be less than message propagation delay
 - In other words, some time it has to be “quite” accurate

Software “Clocks”

- Software “clocks” can incur much lower overhead than maintaining (sufficiently accurate) physical clocks
- Allows a protocol to infer ordering among events
- Goal of software “clocks”: Capture event ordering that are visible to users who do not have physical clocks
 - But what orderings are visible to users without physical clocks?

Visible Ordering to Users without Physical Clocks

- $A \rightarrow B$ (process order)
- $B \rightarrow C$ (send-receive order)
- $A \rightarrow C$ (transitivity)
- $A ? D$
- $B ? D$
- $C ? D$



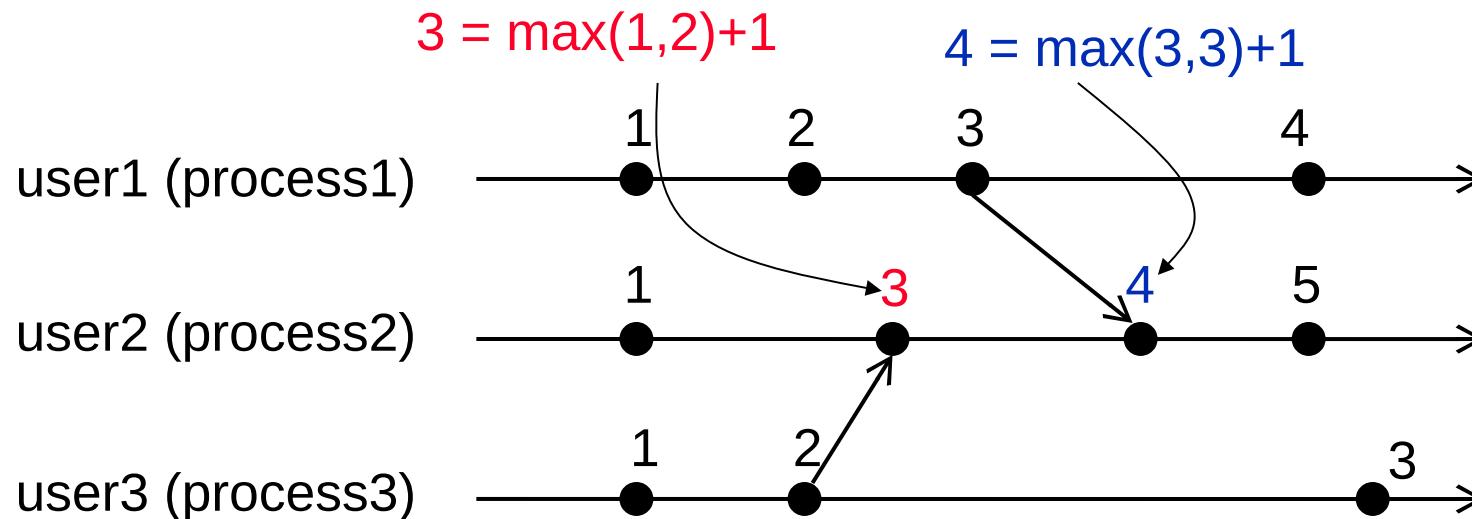
Quick Poll <https://pollev.com/haifengyu229>

“Happened-Before” Relation

- “Happened-before” relation captures the ordering that is visible to users when there is no physical clock
 - A partial order among events
 - Process order, send-receive order, transitivity
- First introduced by Lamport – Considered to be the first fundamental result in distributed computing
 - Lamport won the 1st Edsger W. Dijkstra Prize in Distributed Computing for his work on this
- Goal of software “clock” is to capture the above “happened-before” relation

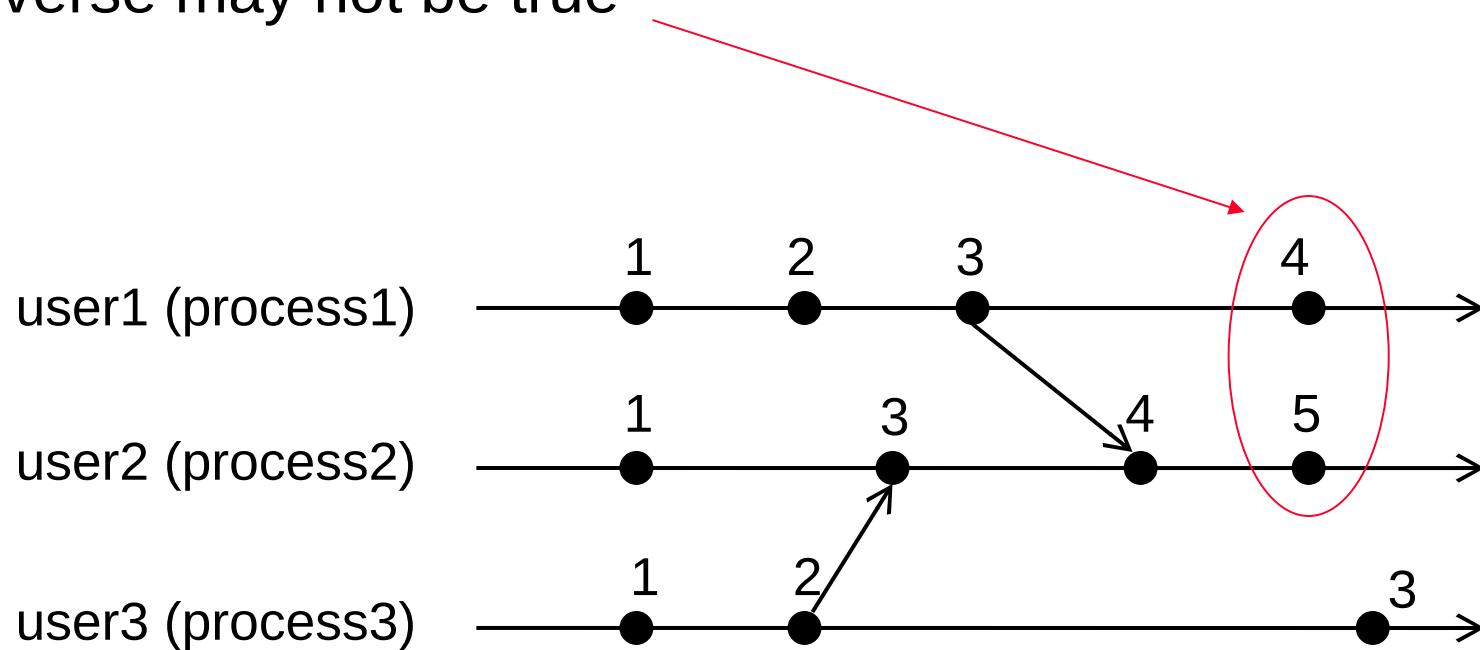
Software “Clock” 1: Logical Clocks

- Each event has a single integer as its logical clock value
 - Each process has a local counter C
 - Increment C at each “local computation” and “send” event
 - When sending a message, logical clock value V is attached to the message. At each “receive” event, $C = \max(C, V) + 1$



Logical Clock Properties

- Theorem:
 - Event s happens before t \Rightarrow the logical clock value of s is smaller than the logical clock value of t.
- The reverse may not be true



Logical Clocks: Deeper Insight

- Assume that logical clock value starts with 1.
- Do the following alternative designs for logical clock work?

Quick Poll <https://pollev.com/haifengyu229>

Software “Clock” 2: Vector Clocks

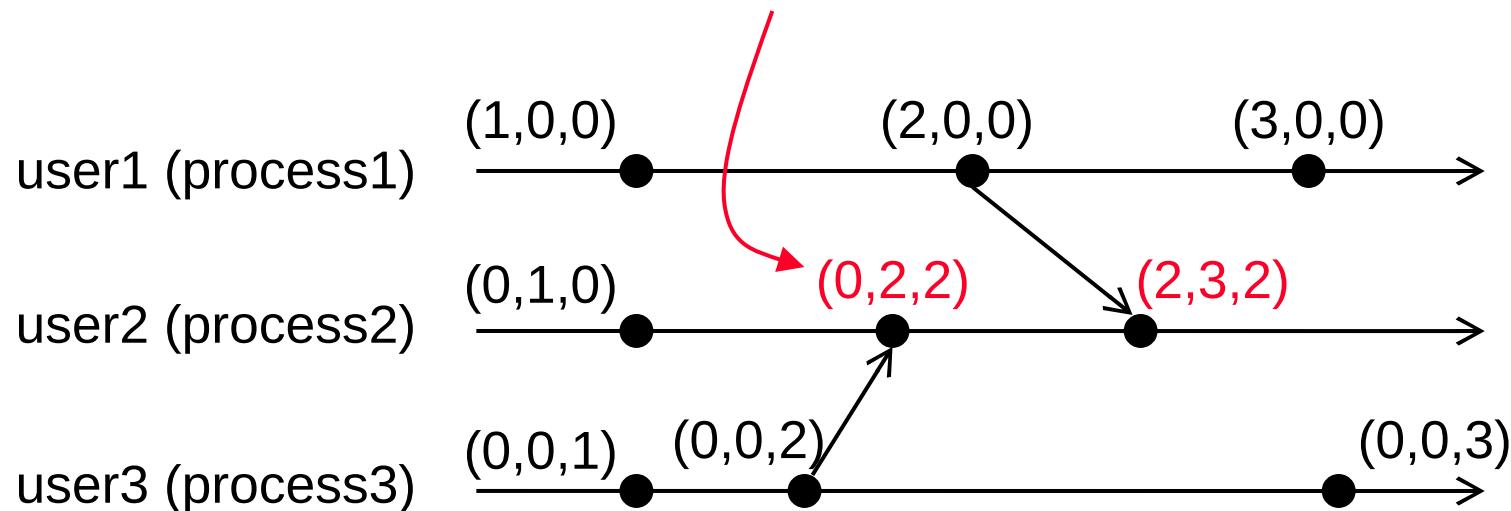
- Logical clock:
 - Event s happens before event t \Rightarrow the logical clock value of s is smaller than the logical clock value of t.
- Vector clock:
 - Event s happens before event t \Leftrightarrow the vector clock value of s is “smaller” than the vector clock value of t.
- Each event has a vector of n integers as its vector clock value
 - $v_1 = v_2$ if all n fields same: $(4,1,7) = (4,1,7)$
 - $v_1 \leq v_2$ if every field in v_1 is less than or equal to the corresponding field in v_2 :
 $(3,1,5) \leq (4,1,7)$
 - $v_1 < v_2$ if $v_1 \leq v_2$ and $v_1 \neq v_2$

Relation “ $<$ here is not
a total order. Example?

Vector Clock Protocol

- Each process i has a local vector C
- Increment $C[i]$ at each “local computation” and “send” event
- When sending a message, vector clock value V is attached to the message. At each “receive” event, $C = \text{pairwise-max}(C, V)$; $C[i]++$;

$$C = (0,1,0), V = (0,0,2)$$
$$\text{pairwise-max}(C, V) = (0,1,2)$$

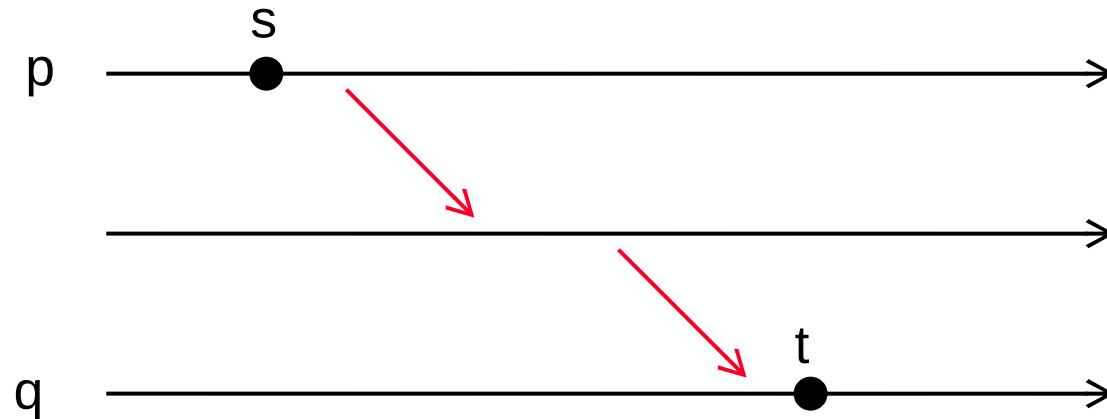


Vector Clock Properties

- Event s happened before t \Rightarrow
vector clock value of s < vector clock value of t.
- Prove by enumeration all possible cases.
- If s happened before t due to process order...
- If s happened before t due to send-receive order...
- If s happened before t due to transitivity, then there must be a chain of events such that
 - s happened before x₁,
 - x₁ happened before x₂,
 - ...
 - x_n happened before t
- Then the vector clock value of s < vector clock value of x₁ < ... < vector clock value of x_n < vector clock value of t.
- Smaller than relation among vector clock values is transitive as well...

Vector Clock Properties

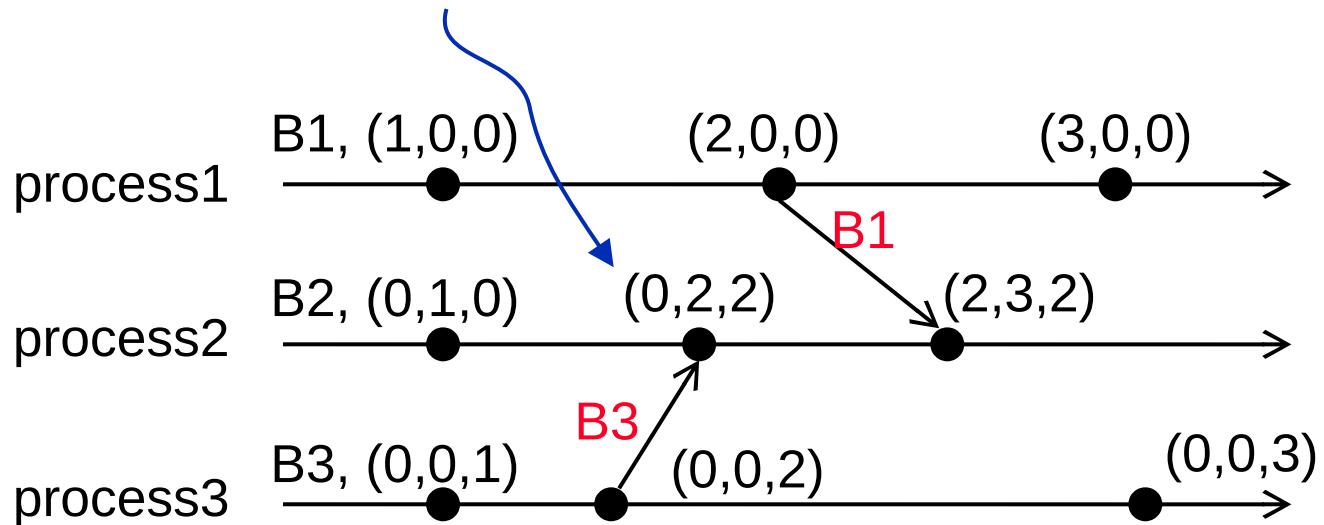
- Vector clock value of $s <$ vector clock value of $t \Rightarrow s$ happened before t
- Prove by separately considering two cases.
- If s and t on same process, done.
- If s is on p and t is on q , let VS be s 's vector clock and VT be t 's.
- $VS < VT \Rightarrow VS[p] \leq VT[p] \Rightarrow$ Must be a sequence of message from p to q after s and before t .



Example Application of Vector Clock

- In Ethereum/Bitcoin, each process has some transactions
 - Or “batches of transactions”
 - Want all processes to know all transactions
- Each transaction has a vector clock value

Process2 has seen all transactions whose vector clock is smaller than $(0,2,2)$.
Based on $(0,2,2)$, process1 knows that process2 has not seen B1.



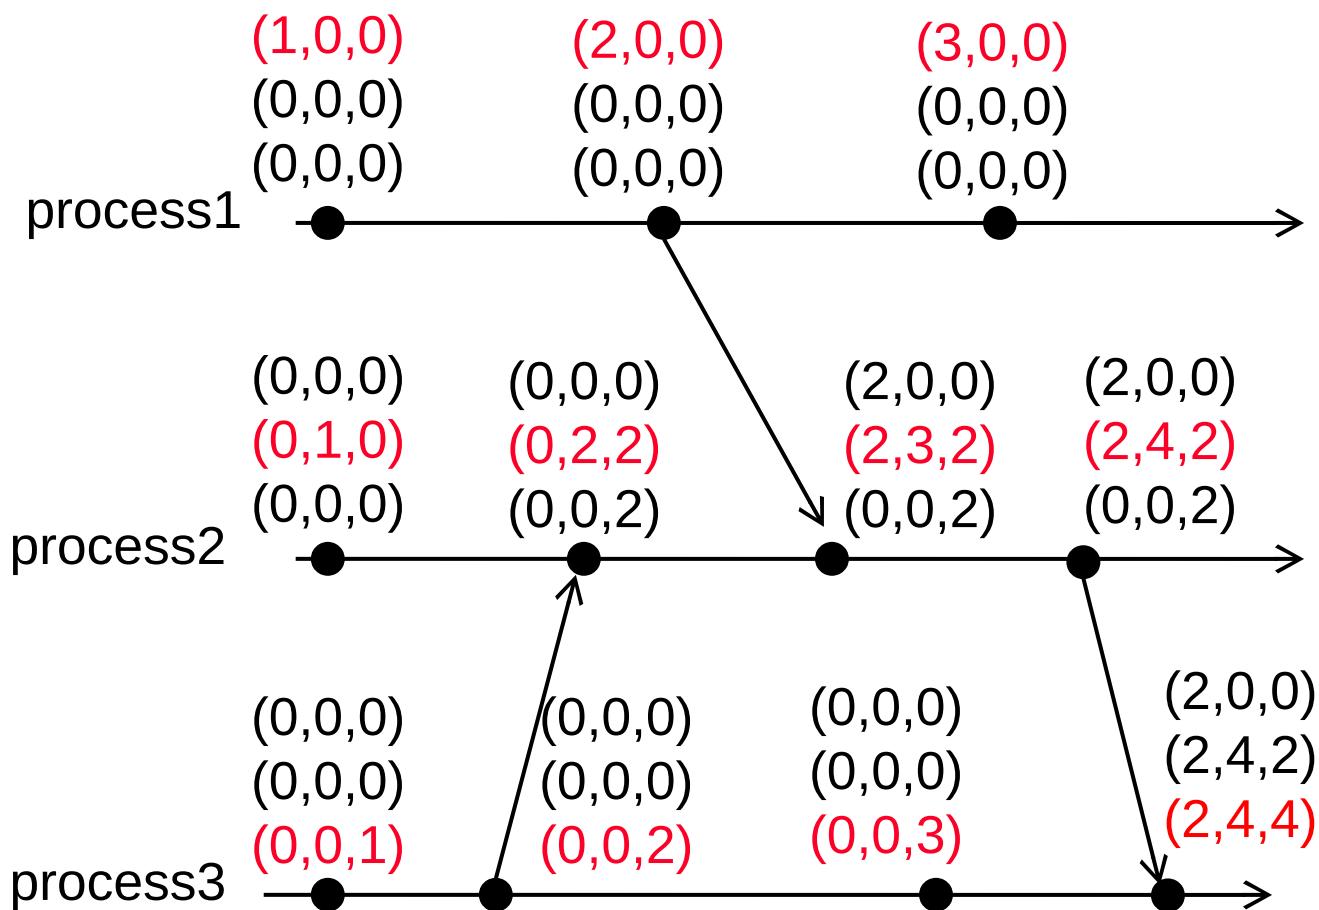
Software “Clock” 3: Matrix Clocks

- Motivation
 - My vector clock describe what I “see”
 - In some applications, I also want to know what other people see
- Matrix clock:
 - Each event has n vector clocks, one for each process
 - The i th vector on process i is called process i ’s principle vector
 - Principle vector is the same as vector clock before
 - Non-principle vectors are just piggybacked on messages to update “knowledge”

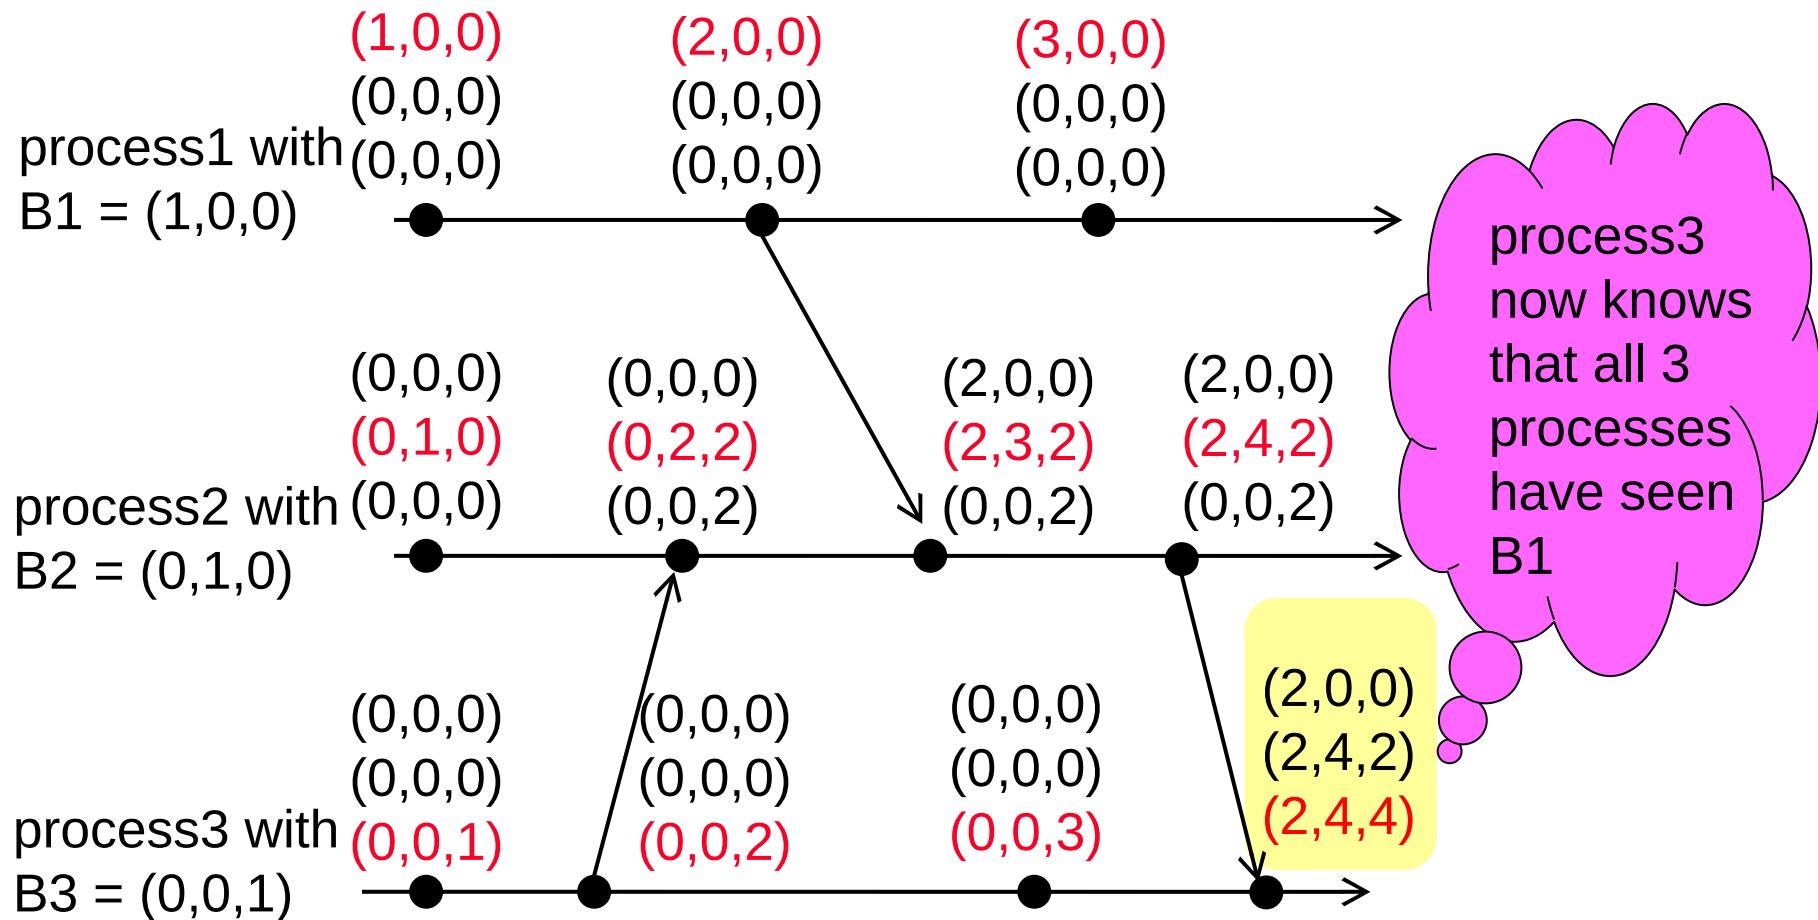
Matrix Clock Protocol

- For principle vector C on process i
 - Increment $C[i]$ at each “local computation” and “send” event
 - When sending a message, all n vectors are attached to the message
 - At each “receive” event, let V be the principle vector of the sender. $C = \text{pairwise-max}(C, V)$; $C[i]++$;
- For non-principle vector C on process i, suppose it corresponds to process j
 - At each “receive” event, let V be the vector corresponding to process j as in the received message. $C = \text{pairwise-max}(C, V)$

Matrix Clock Example



Application of Matrix Clock: Truncated Blockchains



A Snack for Mind

- Vector clock tells me what I know
 - One-dimensional data structure
- Matrix clock tells me what I know about what other people know
 - Two-dimensional data structure
- ?? tells me what I know about what other people know about what other people know
 - ??-dimensional data structure
 - Maybe even ∞ -dimension ?
- Knowledge and common knowledge in distributed systems...

Summary

- “Models and Clocks”
 - Goal: Define “time” in a distributed system
- Logical clock
 - “happened before” \Rightarrow smaller clock value
- Vector clock
 - “happened before” \Leftrightarrow smaller clock value
- Matrix clock
 - Gives a process knowledge about what other processes know

Homework Assignment

- Show that “concurrent with” is not a transitive relation
- We discussed a method by which we can totally order all events within a system. If two events have the same logical time, we broke the tie using process identifiers. This scheme always favor processes with smaller identifiers. Suggest a scheme that does not have this disadvantage. (Hint: Use the value of the logical clock in determining the priority.)
- Assume that you have implemented the vector clock algorithm. However, some application needs Lamport’s logical clock. Write a function *convert* that takes as input a vector timestamp and outputs a logical clock timestamp.
- Bring completed homework to class next week.