

CS4231  
Parallel and Distributed Algorithms

Lecture 6

Instructor: YU Haifeng


## Review of Last Lecture

- “Global Snapshot”
- What is our goal?
- Formalizing consistent global snapshot
- Protocol for capturing a consistent global snapshot

## Today's Roadmap

- “Message Ordering”
- FIFO ordering
  - Already discussed last lecture
- Causal ordering and its application
- Total ordering and its application


## A Internet Chat Room Example



# Formalizing The Notion of Causal Order

- If a “send” event s1 **caused** a “send” event s2
  - **Causal order** requires the corresponding receive event r1 be before r2 (r1 and r2 on the same process)
  - But how do we know whether s1 caused s2?
- If a “send” event s1 **happened before** a “send” event s2
  - Then s1 **may** have caused s2
  - Let’s be pessimistic and be safe -- Assume s1 indeed caused s2
  - Can we avoid being pessimistic?
- **Causal order:** If s1 happened before s2, and r1 and r2 are on the same process, then r1 must be before r2

# How to Ensure Causal Ordering – Intuition



How do you think of this protocol? Is it correct? Any issues?

Quick Poll <https://pollev.com/haifengyu229>


## How to Ensure Causal Ordering – Protocol

- Each process maintains  $n$  by  $n$  matrix  $M$ 
  - This is not the matrix clock!
  - $M[i, j]$ : # of messages sent from  $i$  to  $j$ , as known by process  $i$
- If process  $i$  send a message to process  $j$ 
  - On process  $i$ :  $M[i, j]++$ ;
  - Piggyback  $M$  on the message

## How to Ensure Causal Ordering – Protocol

- Upon process  $j$  receiving a message from process  $i$  with matrix  $T$  piggybacked
  - Let  $M$  be the local matrix on process  $j$
- Deliver the message and set  $M = \text{pairwise-max}(M, T)$ , if
$$\begin{cases} T[k, j] \leq M[k, j] & \text{for all } k \neq i \\ T[i, j] = M[i, j] + 1 \end{cases}$$
  - Intuitively,  $M[i, j]$  on process  $j$  takes consecutive values
  - Otherwise delay message
  - Intuition: *The matrix summarizes the information in the message log in the previous protocol on Slide 6*

## Example Run for the Protocol



## Correctness Proof of the Protocol


- $s_1, s_2, r_1, r_2$ , where  $s_1$  happened before  $s_2$ 
  - Need to prove: both  $r_1$  and  $r_2$  will eventually be delivered – so either  $r_1$  before  $r_2$  or  $r_2$  before  $r_1$
  - Need to prove:  $r_2$  is not before  $r_1$  (we prove this first)
- Consider the case where  $s_1$  and  $s_2$  are on different processes
  - If they are on the same process, it is easier – an exercise for you
- Assume  $s_1$  on process1,  $s_2$  on process2,  $r_1$  and  $r_2$  on process3
- Focus on the top 3 elements of the 3rd column of the matrices

(0,0,0)


(0,0,0)

(0,0,0)


## Correctness Proof



## Correctness Proof



## Correctness Proof



## Correctness Proof (continued)

- Prove by contradiction: Assume Red delivered before Blue
  - After delivering Red: 3rd column of matrix M on process3 will be

$$\begin{array}{r} \geq x_4 \geq x_1 \\ M = \geq y_4 > y_1 \\ 0 \end{array}$$

- M never decreases: Blue can never be delivered any more


To deliver, needs to at least have

$$\underline{M[1, 3] = x_1 - 1}$$

$$M[2, 3] \geq y_1$$

$$M[3, 3] \geq 0$$

The underlined condition  
can never be met



## Need More Correctness Proof

- What we proved so far:
  - If  $s_1$  happened before  $s_2$ , then  $r_2$  will not be before  $r_1$
  - But we don't know if  $r_1$  and  $r_2$  will be delivered at all
- We will now prove that
  - At any given time, there must be one message that can be delivered to the process
    - induction will take care of the other messages

## More Correctness Proof

- Consider a given receiver process  $j$  and its corresponding column in the matrix  $M$   
 $x_1$   
...  
 $x_i$   
...  
 $x_n$
- Consider the non-empty set of all undelivered messages
  - Consider the senders of these messages
  - For each sender  $i$  in the set, there exists an undelivered message whose matrix  $T$  has the column  
 $x_i + 1$   
...

call such undelivered message successor messages

## More Correctness Proof

- Consider the set of successor messages
  - This set must be non-empty
  - This set has at most one message from any given sender
  - This set has at most  $n$  messages total
- Consider all the corresponding “send” events of the successor messages
  - There must be a send event  $s$  that does not have any other send events that happened before  $s$  – why?
  - There can be multiple such  $s$  events
  - Call the corresponding message **top successor message**

## More Correctness Proof

- Claim: Any top successor message can be delivered
  - W.l.o.g, assume the message's sender is process1

$$M = \begin{matrix} x_1 \\ x_2 \\ \dots \\ x_n \end{matrix} \quad T = \begin{matrix} x_1 + 1 \\ y_2 \\ \dots \\ y_n \end{matrix}$$


The matrix column  
on the receiver  
(process  $j$ )      The matrix column  
in the message

need to show:  
 $x_i \geq y_i$  for  $2 \leq i \leq n$

We prove the case for  
 $i = 2$ . Other cases are  
the same.


## More Correctness Proof

- Prove by contradiction: Suppose  $x_2 < y_2$ . Note that this entry corresponds to the # messages sent from process 2 to the receiver process




## More Correctness Proof

- Prove by contradiction: Suppose  $x_2 < y_2$ . Note that this entry corresponds to the # messages sent from process 2 to the receiver process




## More Correctness Proof

- Prove by contradiction: Suppose  $x_2 < y_2$ . Note that this entry corresponds to the # messages sent from process 2 to the receiver process



## More Correctness Proof

- Prove by contradiction: Suppose  $x_2 < y_2$ . Note that this entry corresponds to the # messages sent from process 2 to the receiver process



## Summary of Correctness Proof

- All messages will eventually be delivered
- The delivery order satisfies causal order

## Causal Ordering of Broadcast Messages

- Broadcast: Every message is sent to all people (including the sender itself)
  - Modeled as  $n$  point-to-point messages
- Application: Internet chat room
- For same reason as before, we may need to ensure causal order among messages
  - Each process uses the previous protocol

## Total Ordering of Broadcast Messages

- All messages delivered to all processes in exactly the same order
  - Also called **atomic broadcast**
- Total ordering **only** applies to broadcast messages
- Now assume that all messages in the system are broadcast messages (i.e., every message is sent to all nodes)
  - What is the relation between total order and causal order?

Quick Poll <https://pollev.com/haifengyu229>

## Application: Internet Chat Room

- We want to assign numbers to people have said
  - The number have to be consistent across all users

[0] Alice: Welcome!

[1] Bob: Hello, Alice

[2] Alice: Let's try to prove P=NP

[3] Bob: OK, where do we start?

....

....

....

[1712] Bob: I am confused, what were we trying to prove?

[1713] Alice: Please refer to message [2]


## Application: Distributed Ledger / Blockchain

- Each node maintains a ledger of transactions
  - Alternatively, each node maintains a list of blocks, where each block contains a list of transactions
- Each new transaction is broadcast to all nodes
- We want all the nodes to have the same ordering for all the transactions
  - This ensures consistency across the ledgers on all nodes
- Our protocol next are assuming no failures...

## Using a Coordinator for Total Order Broadcast


- A special process is assigned as the **coordinator**
- To broadcast a message
  - Send a message to the coordinator
  - Coordinator assigns a sequence number to the message
  - Coordinator forward the message to all processes with the sequence number
  - Messages delivered according to sequence number order
- Problem:
  - Coordinator has too much control

## Skeen's Algorithm for Total Order Broadcast




- Each process maintains
  - Logical clock and a message buffer for undelivered messages
- A message in the buffer is delivered / removed if
  - All messages in the buffer have been assigned numbers
  - This message has the smallest number

## Correctness Proof for Skeen's Algorithm



- Claim: All messages will be assigned message numbers
- Claim: All messages will be delivered
- Claim: If message A has a number smaller than B, then B is delivered after A – Prove by contradiction on next slide (trivial?)

## Correctness Proof for Skeen's Algorithm -- Continued



key: Process 3's logical clock now must be larger than B's number

- Suppose A is delivered on process 3 after B.
- Then A must have been placed in buffer after B was delivered
- A must have a number larger than B – Contradiction.

# Summary

- “Message Ordering”
- FIFO ordering for point-to-point messages
  - Already discussed last lecture
- Causal ordering for point-to-point messages
  - Applications
  - Protocol to ensure causal ordering
- Causal ordering for broadcast messages
  - Protocol
- Total ordering for broadcast messages
  - Application
  - Skeen’s algorithm

## Homework Assignment (on this and next few slides)

- Show the relationship between conditions (C1), (C2), and (C3) on message delivery of a system.

$$(C1): \quad s_1 \rightarrow s_2 \Rightarrow \neg(r_2 \rightarrow r_1)$$

$$(C2): \quad s_1 < s_2 \Rightarrow \neg(r_2 \rightarrow r_1)$$

$$(C3): \quad s_1 \rightarrow s_2 \Rightarrow \neg(r_2 < r_1)$$

Here  $s_1$  and  $s_2$  are sends of any two messages, and  $r_1$  and  $r_2$  are the corresponding receives. Note that a computation satisfies a delivery condition if and only if the condition is true for all pairs of messages. Here " $s_1 \rightarrow s_2$ " means that  $s_1$  **happened-before**  $s_2$ , and " $s_1 < s_2$ " means that  $s_1$  and  $s_2$  are **on the same process** and  $s_1$  is **before**  $s_2$ .

For each pair  $(x, y)$  of the three conditions (i.e., C1 & C2, C1 & C3, C2 & C3)

- Prove  $x$  implies  $y$  and  $y$  implies  $x$ , OR
- Prove  $x$  implies  $y$ , and a counter example that satisfies  $y$  but not  $x$ , OR
- Prove  $y$  implies  $x$ , and a counter example that satisfies  $x$  but not  $y$ , OR
- A counter example that satisfies  $y$  but not  $x$ , and a second counter example that satisfies  $x$  but not  $y$

- For the protocol on Slide 23, does the resulting total order always satisfy causal order? Why?
- Bring your completed homework to class next week.